REMOTE PROCEDURE CALL

51

51

RPC as a Programming Abstraction

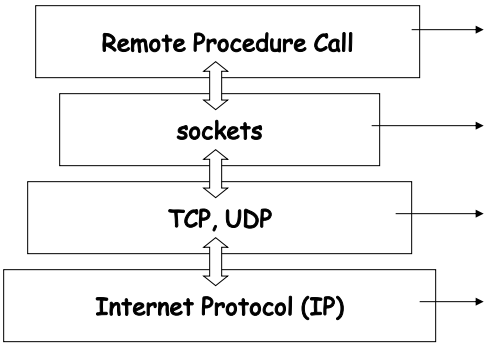

Remote Procedure Call:
hides communication details behind a procedure call and helps bridge heterogeneous platforms

sockets:
operating system level interface to the underlying communication protocols

TCP, UDP:
User Datagram Protocol (UDP) transports data packets without guarantees
Transmission Control Protocol (TCP) verifies correct delivery of data streams

Internet Protocol (IP):
moves a packet of data from one node to another

52

52

The basic RPC protocol

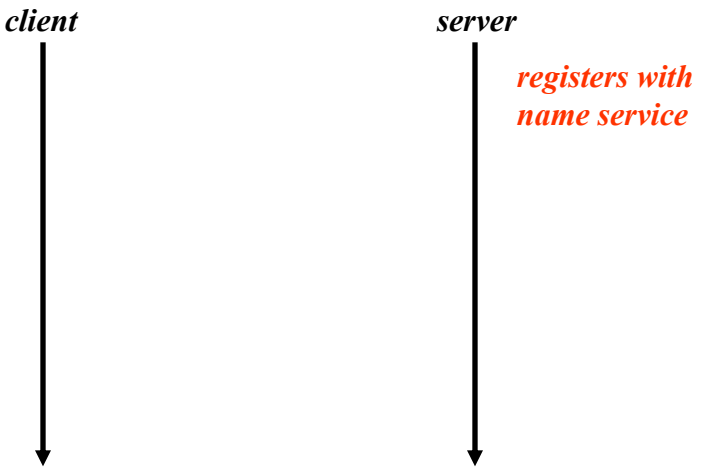

53

53

The basic RPC protocol

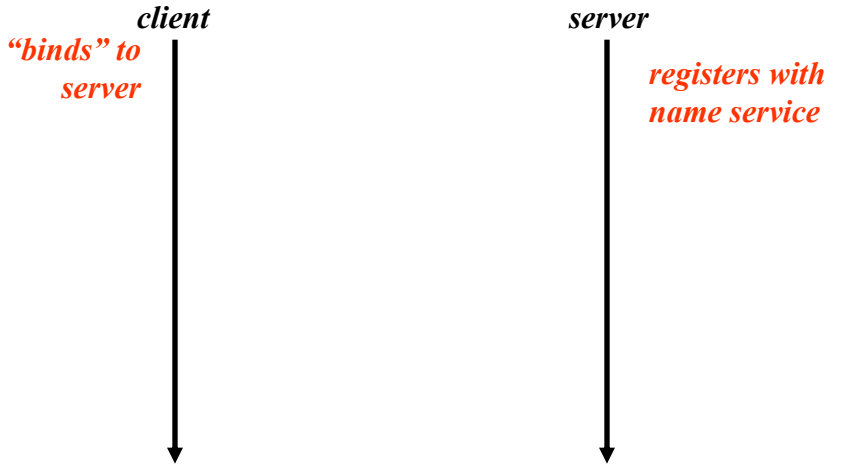

54

54

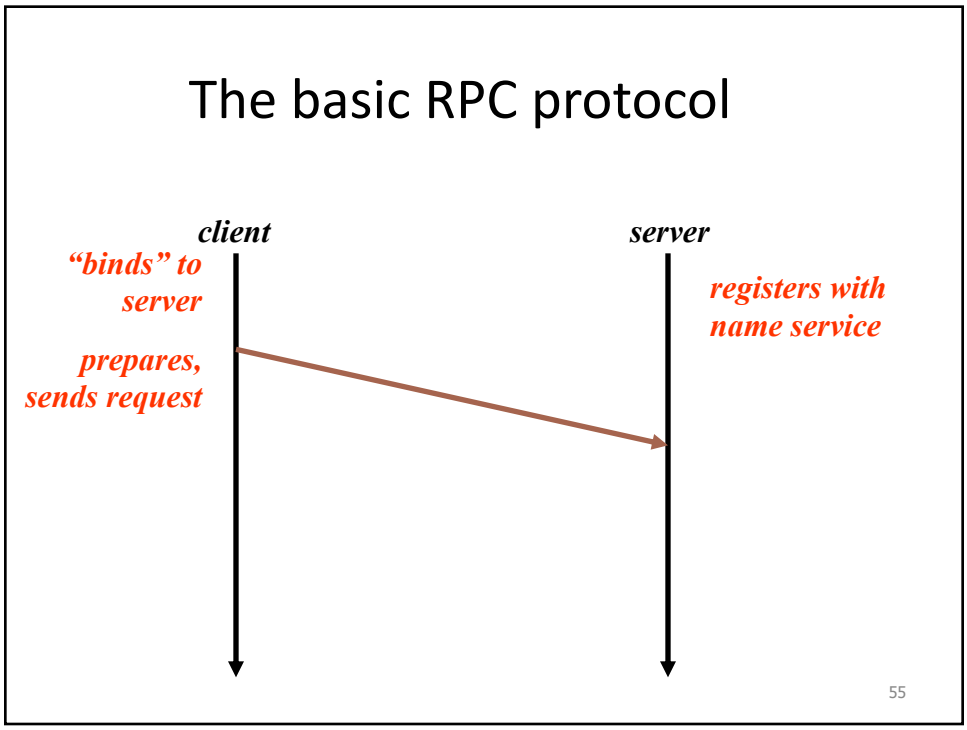
The basic RPC protocol
client
server
"binds" to server
registers with name service
prepares, sends request
55

55

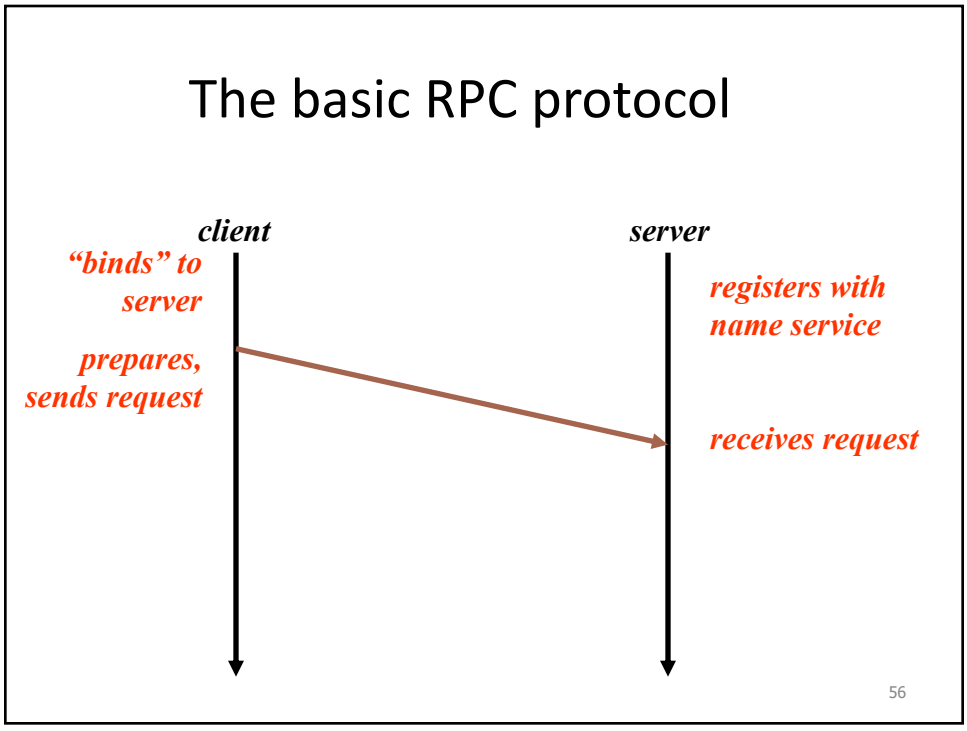
The basic RPC protocol
client
server
"binds" to server
registers with name service
prepares, sends request
receives request
56

56

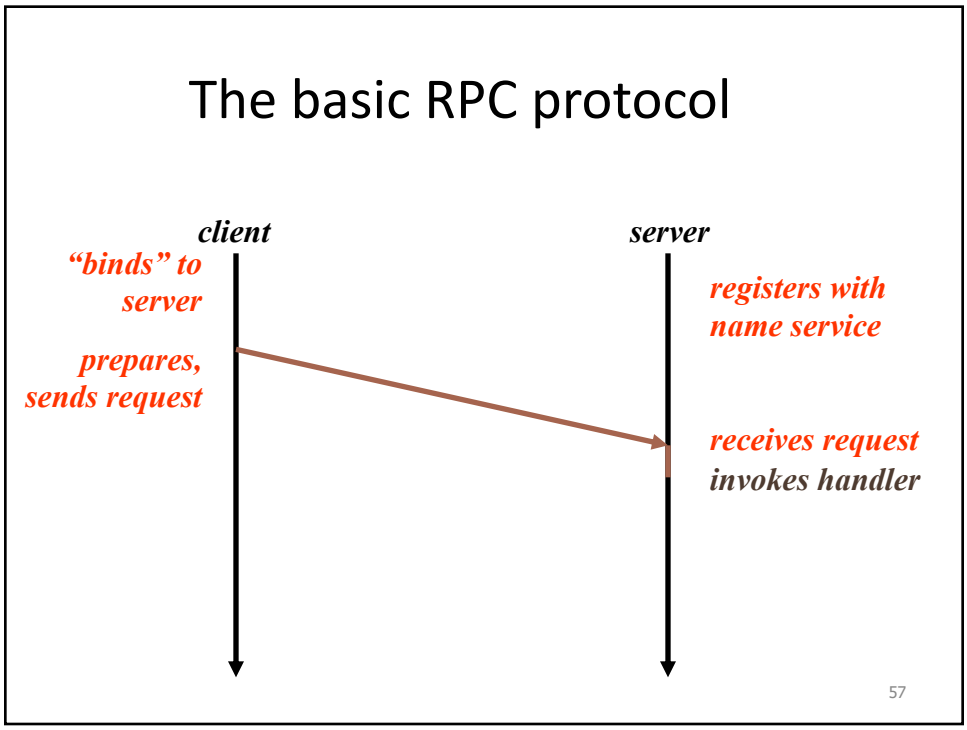
The basic RPC protocol
client
server
"binds" to server
registers with name service
prepares, sends request
receives request
invokes handler
57

57

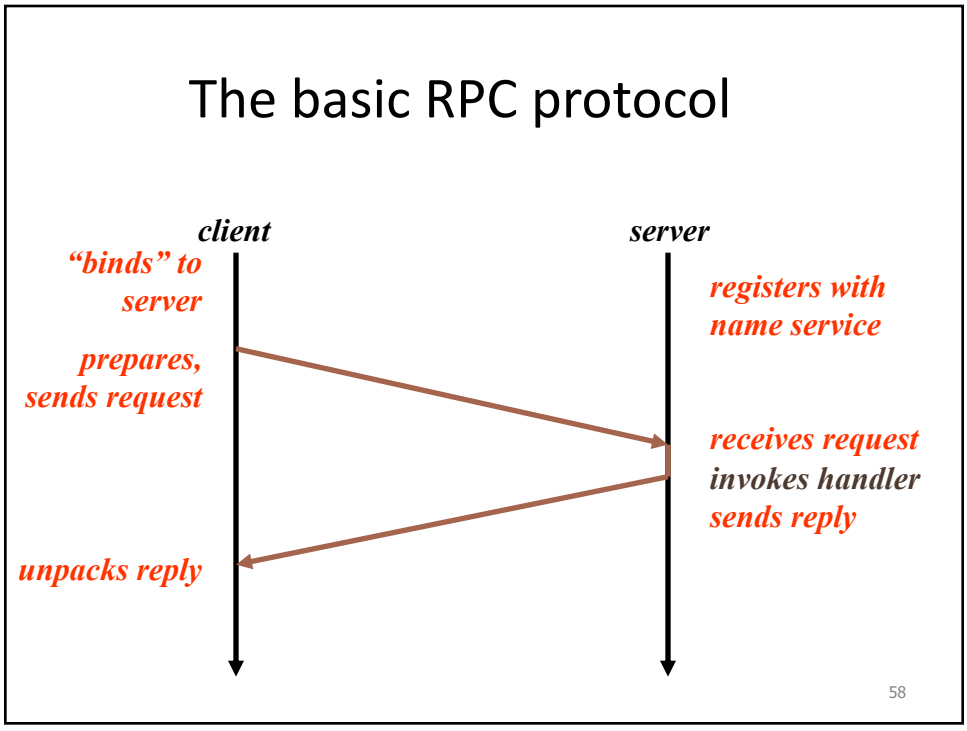
The basic RPC protocol
client
server
"binds" to server
registers with name service
prepares, sends request
receives request
invokes handler
sends reply
unpacks reply
58

58

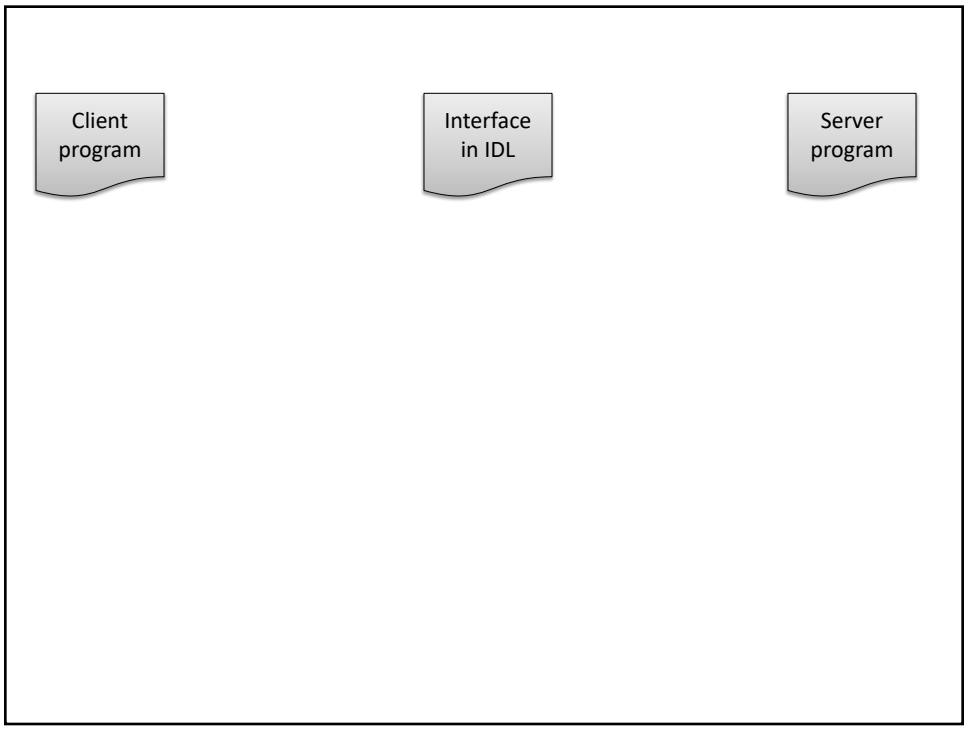
Client program
Interface in IDL
Server program

59

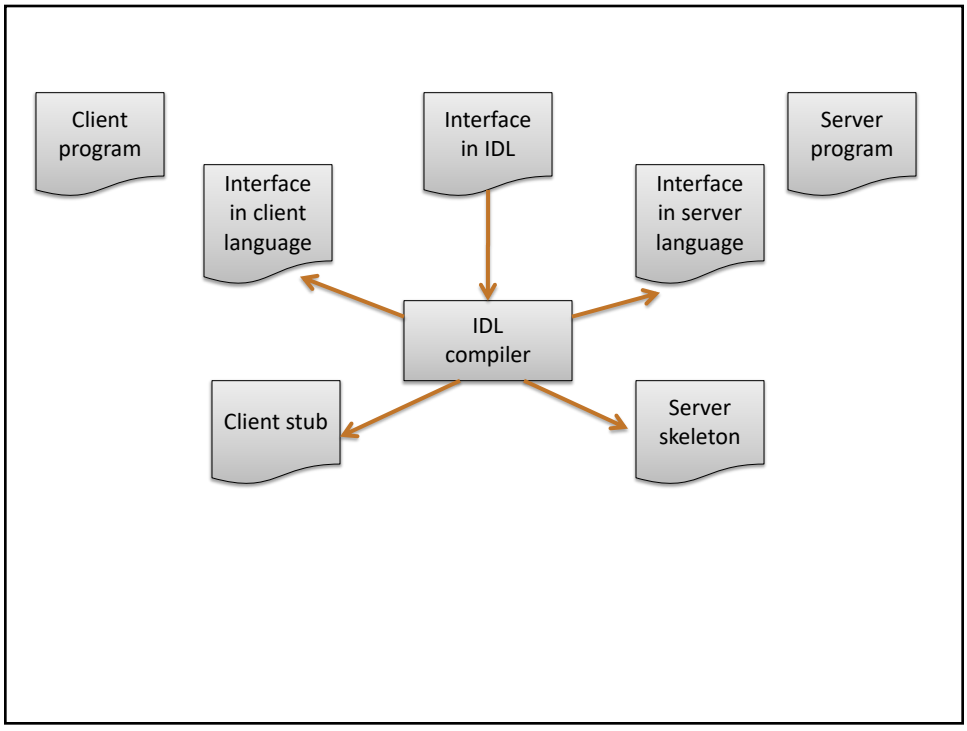
Client program
Interface in IDL
Server program
Interface in client language
Interface in server language
IDL compiler
Client stub
Server skeleton

60

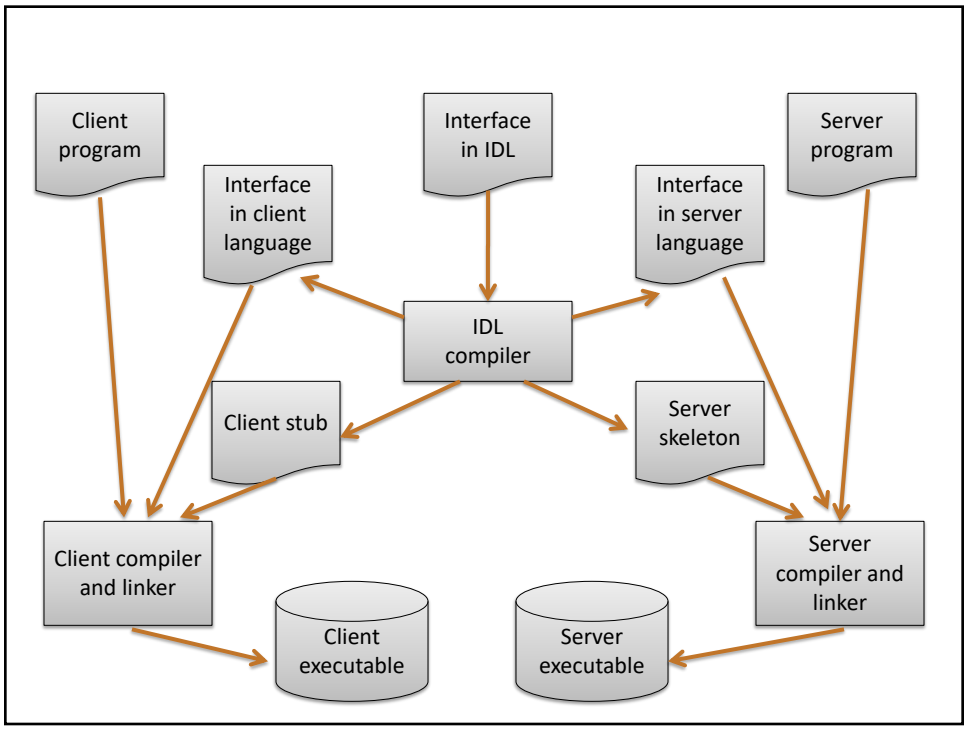

61

RPC Binding

62

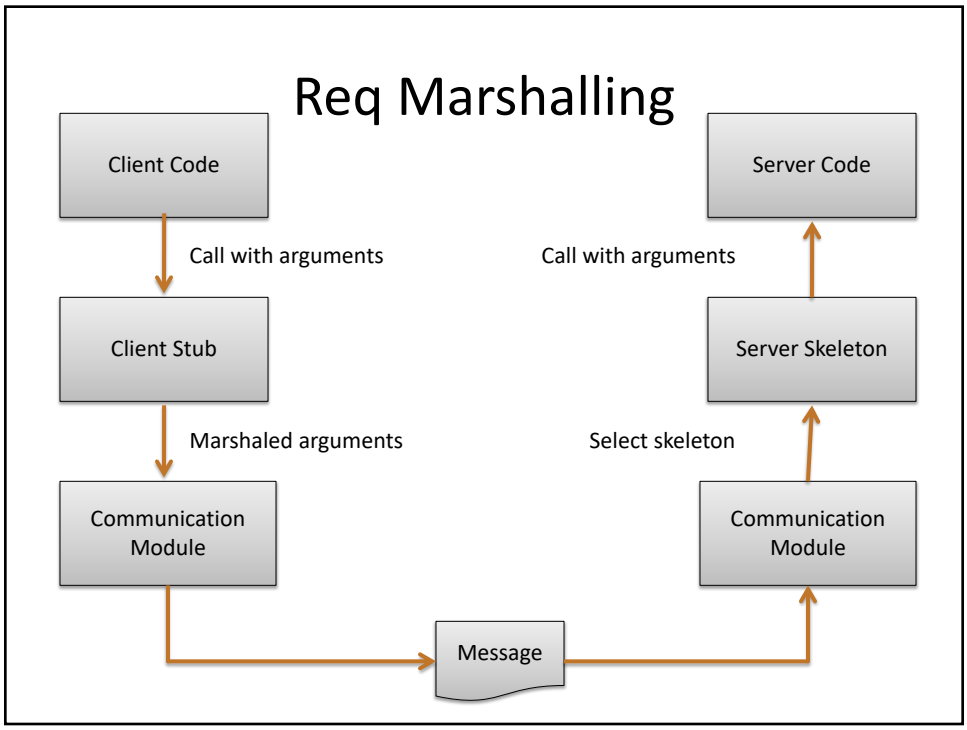

63

Data in messages

- Marshalling
- Byte ordering issues (big-endian versus little endian), strings (some CPUs require padding), alignment, etc
- As fast as possible
  - yet must also be very general

64

64

Fancy argument passing

- RPC transparent for simple calls
  - New forms of exceptions
- Complex structures, pointers, big arrays?
- Most limit size, types of RPC arguments.

65

65

RPC WITH LOST MESSAGES

66

66

RPC Semantics in the Presence of Failures

- The client is unable to locate the server
- The request message from the client to server is lost
- The reply message from the client is lost
- The server crashes after sending a request
- The client crashes after sending a request

Client is Unable to Locate Server

- Causes:
  - server down
  - server moved
  - different version of server, …
- Fixes
  - Network failure exceptions
    - Transparency is lost

Messages Lost

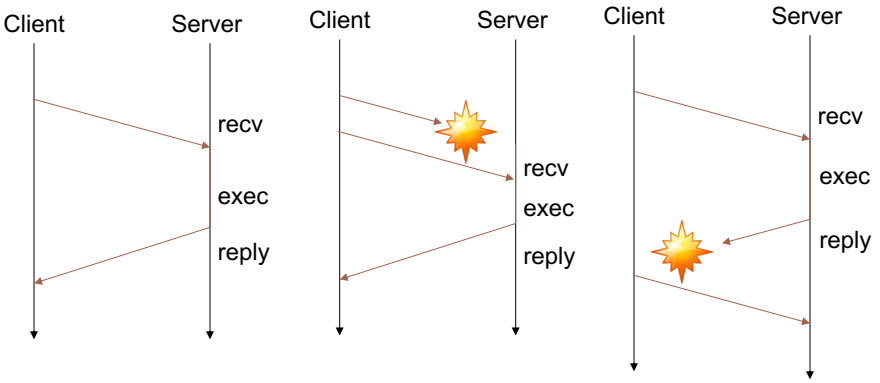

69

69

Lost Request Message

- Easiest to deal with
- Just retransmit the message!
- If multiple message are lost then
  - "client is unable to locate server" error

70

70

Overcoming lost packets

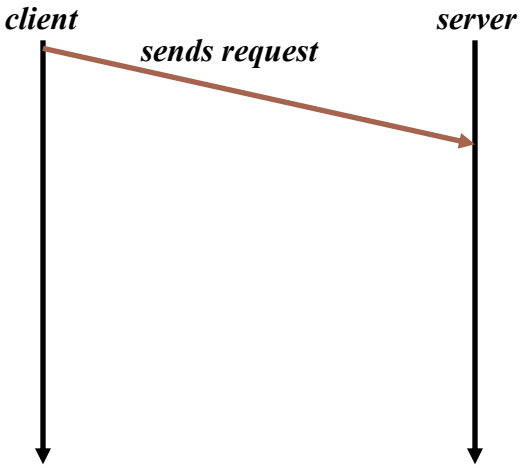

71

71

Overcoming lost packets

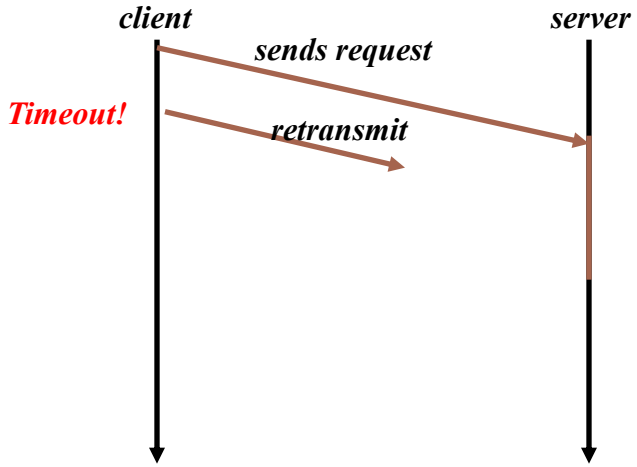

72

72

Overcoming lost packets

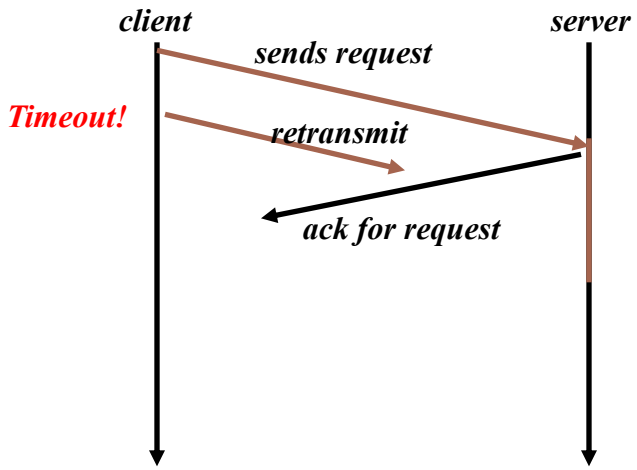

73

73

Overcoming lost packets

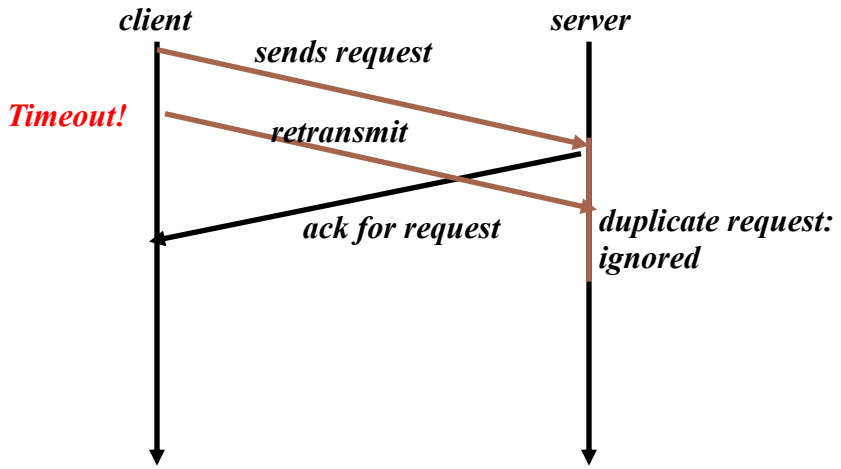

74

74

Overcoming lost packets

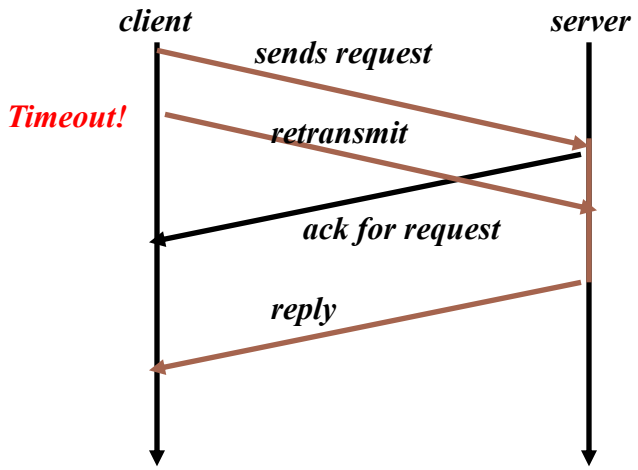

75

75

Overcoming lost packets

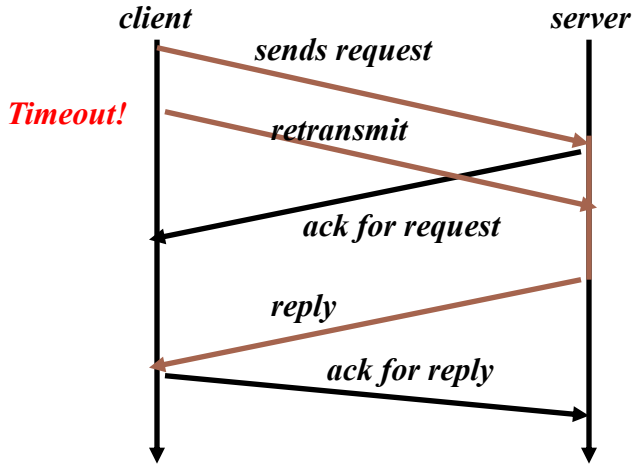

76

76

Costs in fault-tolerant version?

- Acks are expensive.
  - Suppress the initial ack
- Retransmission is costly.
  - Tune the delay
- Big messages
  - ack a burst at a time

77

77

Big packets

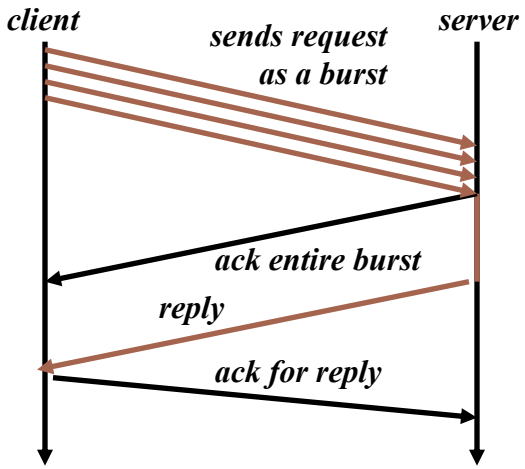

78

78

Lost Reply Message

- Did server execute the procedure or not?
- Possible fixes
  - Retransmit the request
    - Only works if operation is idempotent
  - What if operation not idempotent?
    - Assign unique sequence numbers to every request
    - **Shared state between client and server…**

79

79

RPC WITH CRASHES

80

80

Server Crashes

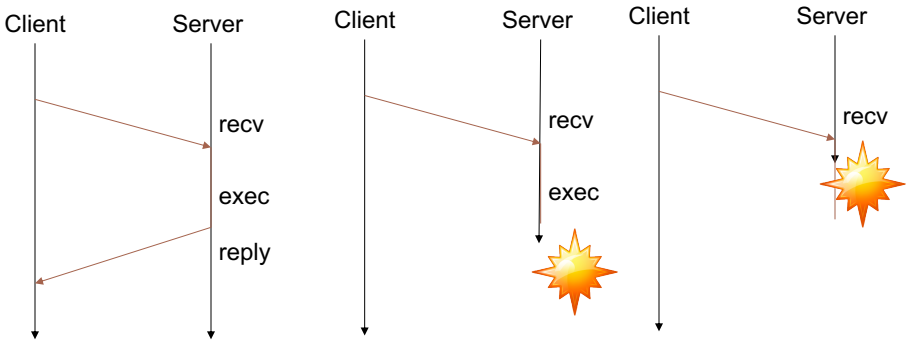

81

81

Server Crashes

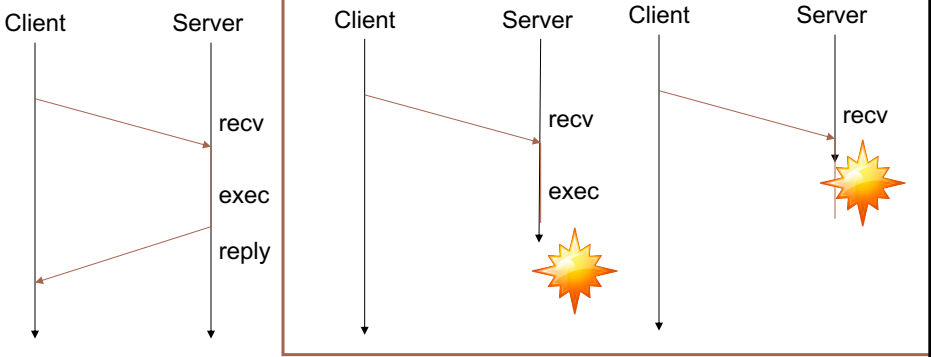

Client cannot distinguish these scenarios

82

82

Server Crashes

- Three possible semantics
  - At least once semantics
    - Client keeps trying until it gets a reply
  - At most once semantics
    - Client gives up on failure
  - Exactly once semantics
    - Can this be correctly implemented?

Impossibility of Exactly Once Semantics (1)

- Example: Print server
  - Client wants to print a document on the server
  - Three possible events at server:
    - Reply with completion message (R)
    - Print the text (P)
    - Crash (C)

Impossibility of Exactly Once Semantics (2)

- These events (R, P, C) can occur in six different orderings:

Reply	Reply	**Crash**	Print	Print	**Crash**
Print	**Crash**	(Reply)	Reply	**Crash**	(Print)
Crash	(Print)	(Print)	**Crash**	(Reply)	(Reply)

85

85

Impossibility of Exactly Once Semantics (3)

- Assume server crashes, later recovers, and notifies clients
- What should clients do?

Client reissue strategy	Server (strategy R → P)		
	RPC	RC(P)	C(RP)
Always	DUP	OK	OK
Never	OK	ZERO	ZERO
Only when ACK	DUP	OK	ZERO
Only when not ACK	OK	ZERO	OK

86

86

Impossibility of Exactly Once Semantics (4)

- Assume server crashes, later recovers, and notifies clients
- What should clients do?

Client reissue strategy	Server (strategy P → R)		
	PRC	PC(R)	C(PR)
Always	DUP	DUP	OK
Never	OK	OK	ZERO
Only when ACK	DUP	OK	ZERO
Only when not ACK	OK	DUP	OK

87

87

Impossibility of Exactly Once Semantics (5)

- Can we fix this?
- Fundamental problem: distributed agreement
- Analogy: Coordinating attacks of allied armies
  - Allied armies on opposing hillsides
  - Messengers travel through valley occupied by enemy
  - Omission failures (not Byzantine!)

88

88

Impossibility of Exactly Once Semantics (6)

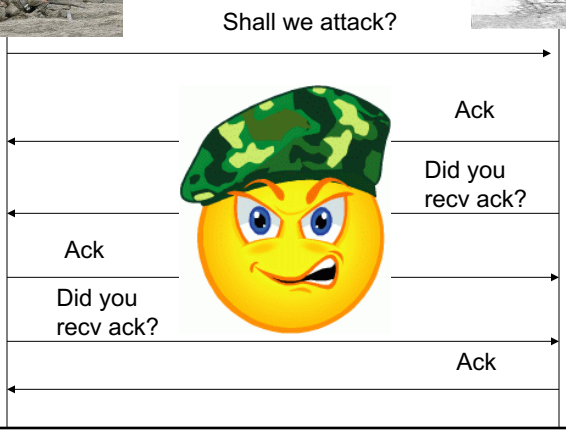

89

89

Client Crashes

- Let's the server computation **orphan**
- Orphans can
  - Waste CPU cycles
  - Lock files
  - Client reboots and it gets the old reply immediately

90

90

Client Crashes: Possible Solutions

- Extermination:
  - Client log, kill orphan on reboot
  - Disadvantage: log overhead
- Reincarnation:
  - Client broadcasts new “epoch” when reboots
  - New epoch: servers kill orphans
  - Disadvantage: network partition

91

91

Client Crashes: Possible Solutions

- Expiration:
  - Each RPC is given a lease T
  - Must renew lease before expiration
  - If client reboots after T sec, all orphans are gone
  - Problem: what is a good value of T?

92

92