LOCAL RPC

RPC versus local procedure call

- Restrictions on argument sizes and types
- New error cases:
  - Bind operation failed
  - Request timed out
  - Argument “too large”
- Costs may be very high

RPC costs in case of local destination process

- Often, the destination is right on the caller's machine
  - Caller builds message
  - Send system call, then receive, blocks, context switch
  - Message copied into kernel, then out to dest.
  - Context switch
  - Dest computes result
  - Repeated in reverse direction
  - If scheduler is a process, may context switch 4 times!

117

117

RPC example

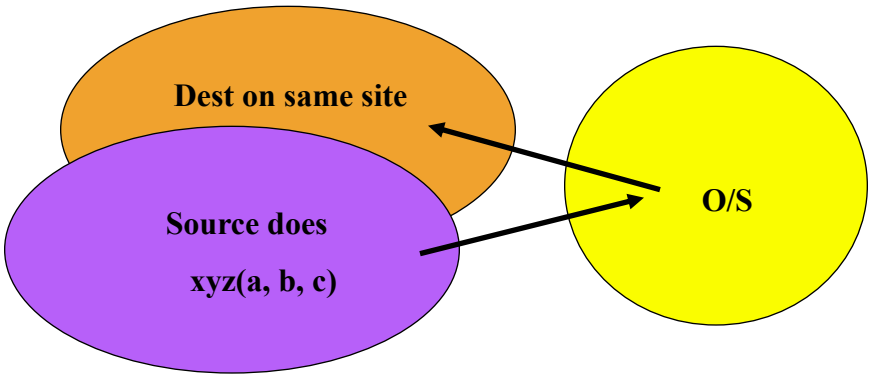

118

118

RPC in normal case

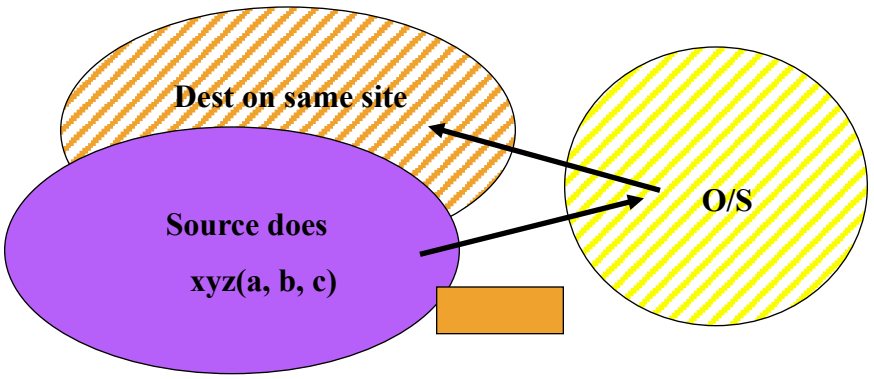

119

119

RPC in normal case

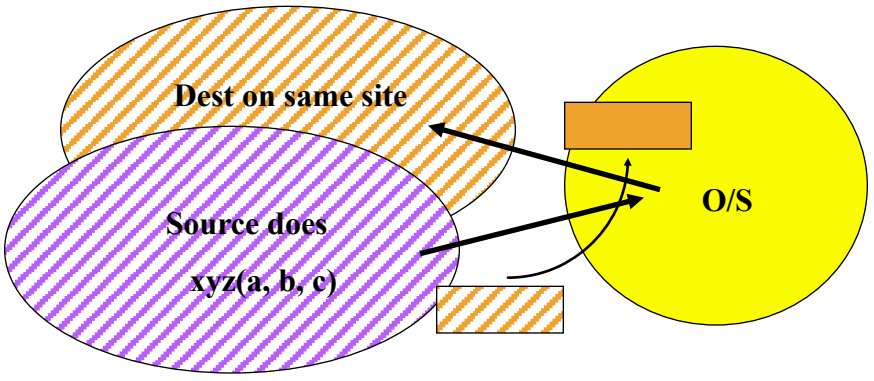

120

120

RPC in normal case

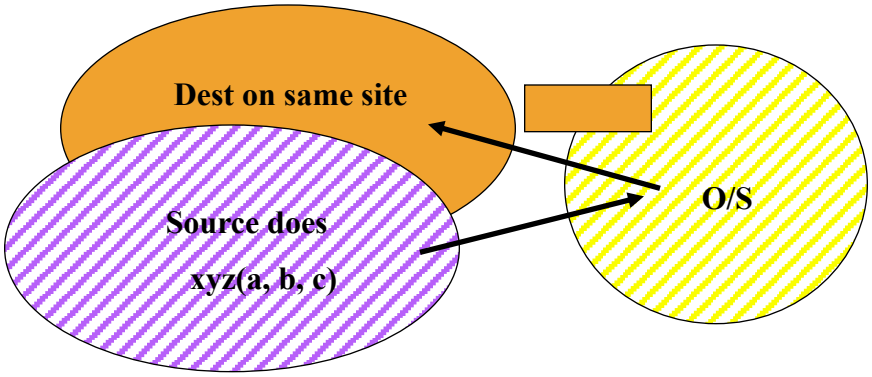

121

121

Important optimizations: LRPC

- Lightweight RPC (LRPC):
- Uses memory mapping to pass data
- Reuse kernel thread
- Single system call: send_rcv or rcv_send

122

122

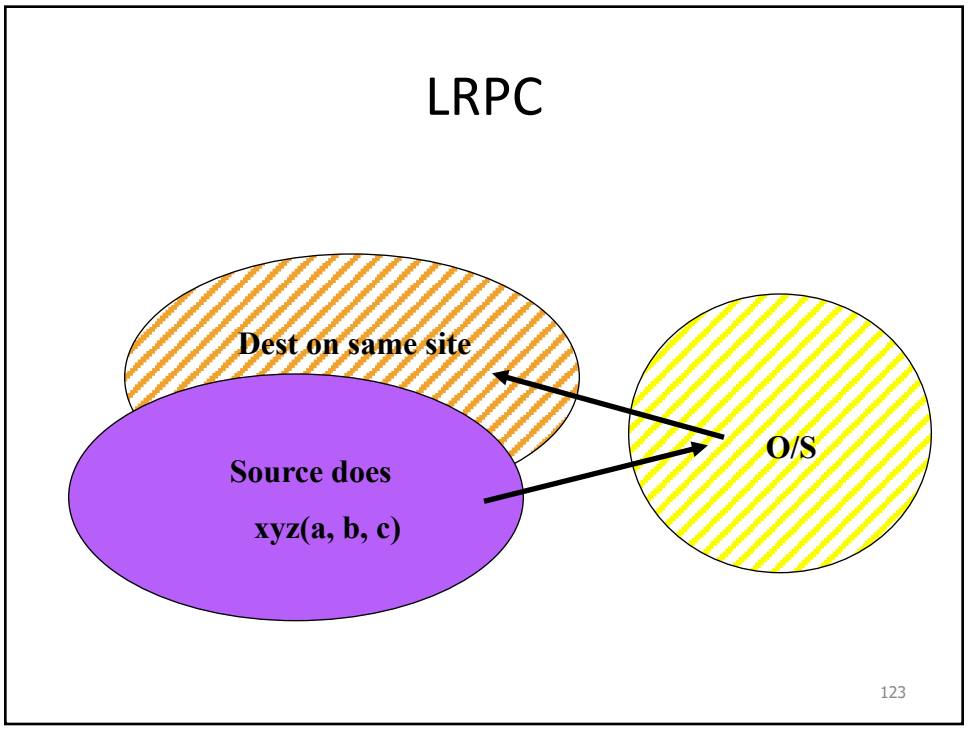
LRPC
Dest on same site
Source does
xyz(a, b, c)
O/S
123

123

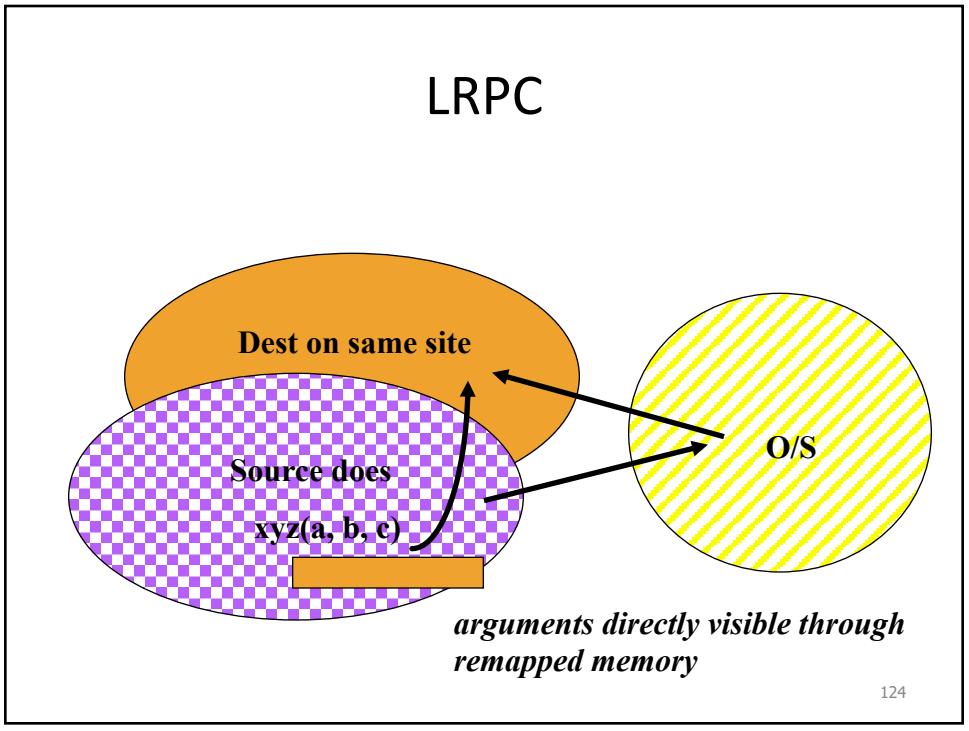
LRPC
Dest on same site
Source does
xyz(a, b, c)
O/S
arguments directly visible through
remapped memory
124

124

LRPC performance impact

- 10-fold improvement over hand-optimized RPC implementation
- Two memory remappings, no context switch
- 50 times faster than standard RPC
- Easy to ensure exactly once

125

125

FBUFS

126

126

Fbufs

- Speed up layered protocols
- Buffer management overhead
- Solution: “cache” buffers
  - memory management & protection
- Stack layers share memory

127

127

Fbufs

control flows through stack of layers, or pipeline of processes

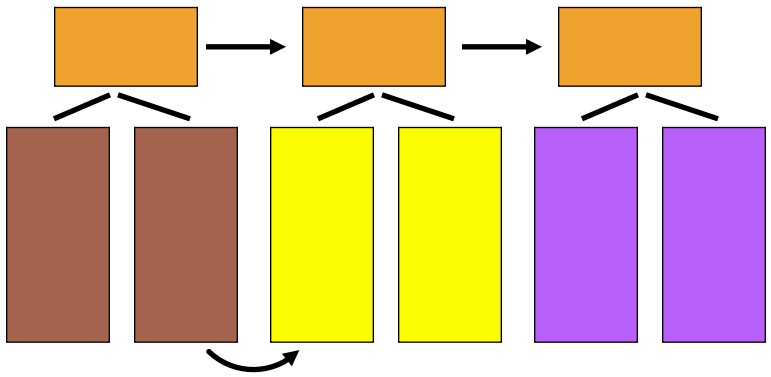

128

128

Fbufs

control flows through stack of layers, or pipeline of processes

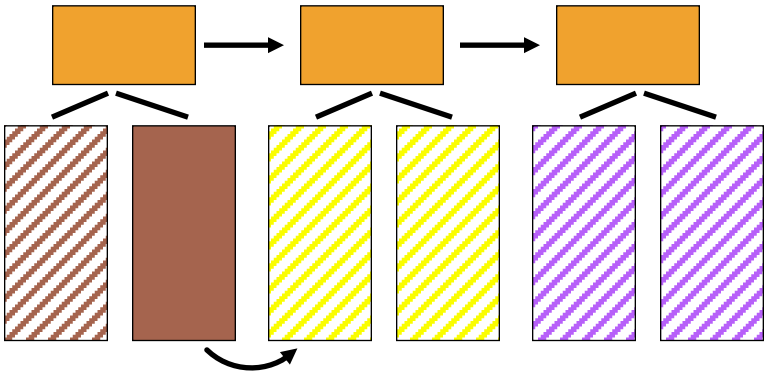

129

129

Fbufs

control flows through stack of layers, or pipeline of processes

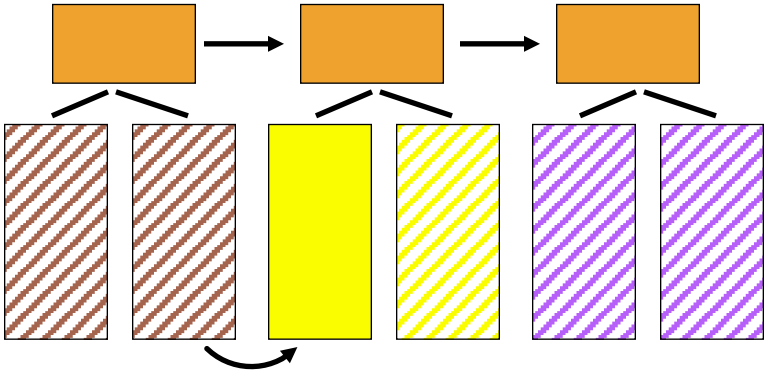

130

130

Fbufs

control flows through stack of layers, or pipeline of processes

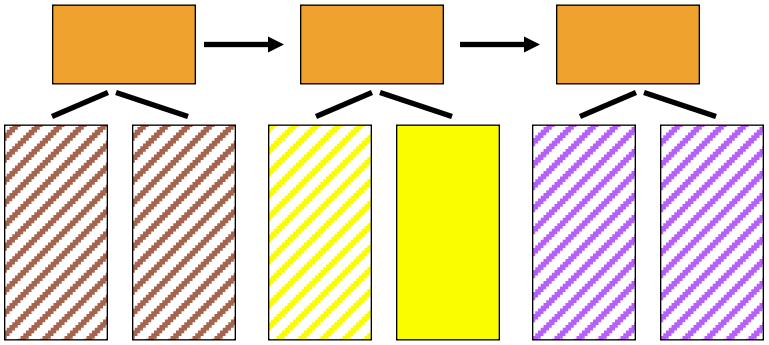

131

131

Fbufs

control flows through stack of layers, or pipeline of processes

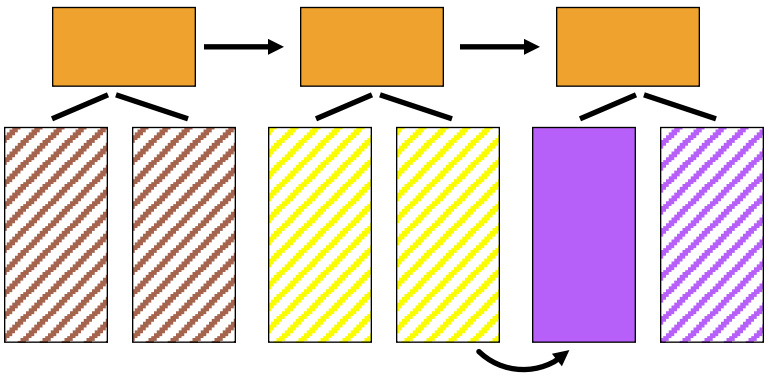

132

132

Where are Fbufs used?

- Most kernels use similar ideas
- Many application-layer libraries

133