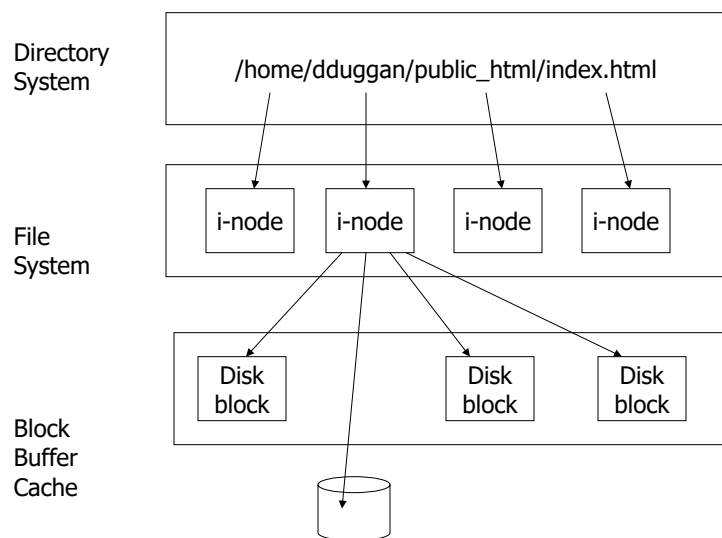


FILE SYSTEMS

2

2

Layers in File Systems





3

3

Naming Scheme

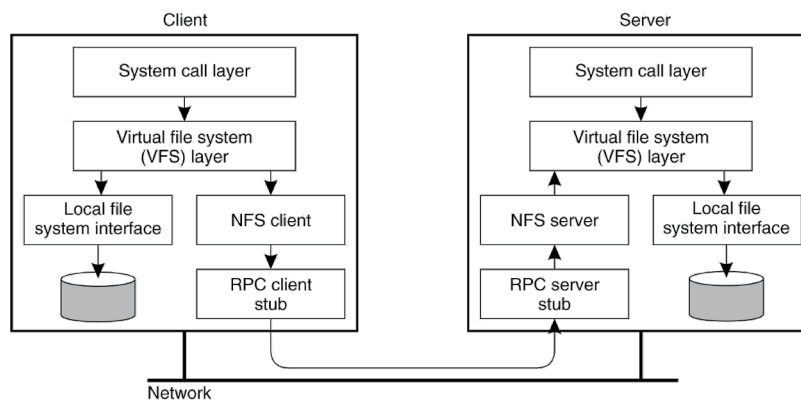
- **Mount remote directories to local**
 - Coherent **local** directory tree
 - Ex: Unix/Linux with NFS; Windows mapped drives
- **Total integration of component file systems**
 - Single **global** name structure
 - Ex: AFS



4

4

Location Transparency




5

5

Location Independence (Migration Transparency)

- NFS: Not location independent
 - Server: `export /root/fs1/`
 - Client: `mount server:/root/fs1 /fs1`




6

6

Location Independence (Migration Transparency)

- AFS: global volumes
 - Global directory /afs;
 - /afs/cs.stevens.edu/vol1/...; /afs/cs.njit.edu/vol1/
 - File id = <vol_id, vnode #, uniquifier>
 - “Volume location database”
 - vol_id → server_ip mappings
 - Shared by servers



7

7

File Server Semantics

- Stateless
 - + Crash recovery
 - - File locking
 - Ex: NSF v3
- Stateful
 - - Crash recovery
 - + File locking
 - Ex: AFS, NFS v4

8

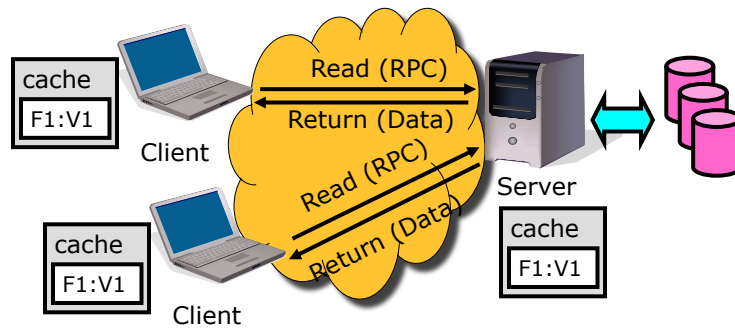
8

CACHING POLICIES

9

9


File Caching



10

10

Cache Consistency



11

11

Cache Update Policies

- When does the client update the master file?
- **Write-through:** write data to server ASAP
- **Delayed-write:** cache, write to server later
 - Better local performance
 - Network I/O
 - Poor reliability
- **Write-on-close**

12

12

File Sharing Semantics

- **Sequential Semantics**
 - No cache
 - Performance problems
 - Write-through cache
 - Must notify clients holding copies
- **Session Semantics**
- **Immutable Files**
- **Atomic Transactions**

13


13

NFS: NETWORK FILE SYSTEM

14

14

Caching in NFS



15

15

Client Caching

- Client checks validity of cached files
 - File open
 - Periodic polling
- Client responsible for writing out cache
 - Periodic scan, flush of dirty blocks

16

16

NFS Semantics

- Locking
 - Originally separate (stateless)
 - Stateful for locking since NFS v4
- Unix file semantics not guaranteed
 - E.g., read after write
- Session semantics not even guaranteed
 - Intermediate writes
 - Client may implement close-to-open

17

17

NFS Implementation

- Remote procedure calls for all operations
 - Originally over UDP
 - Using TCP since NFSv4
- Lost requests are simply re-transmitted
 - At-least-once semantics

18

18


NFS Failure Recovery

- Server crashes transparent to client
 - Each client request self-contained
 - Client retransmits request if crash
 - *“Server not responding, still trying”*
- Client crashes transparent to server

19

19

Caching and Failures

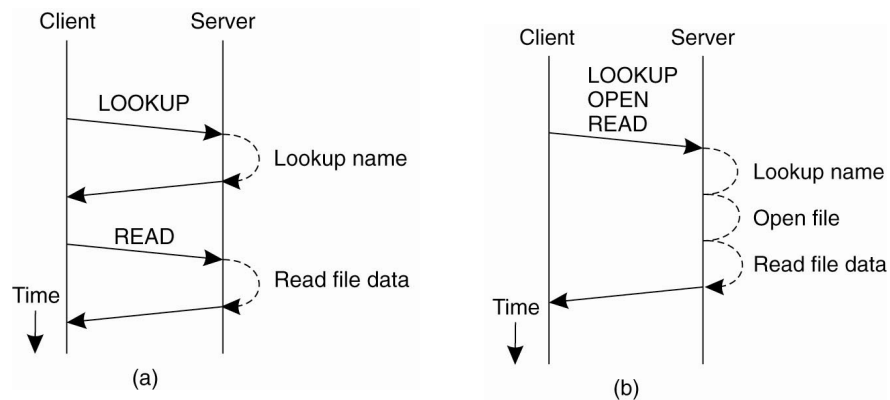


- Suppose client implements close-to-open
- Server acks updates
- Server crashes before flushing updates to disk
- Fixes:
 - Server flushes updates before ack
 - NVRAM for server
 - NFS v3: client buffers updates until COMMIT acknowledged

20

20

Compound RPC



21

21

Open Delegation

- Server may delegate open/close/locking to client
- Operations done locally at client
- Periodic cache checks unnecessary
- Lease and revocation via callback
 - RPC from server to client

22

22