

Distributed Hash Tables II

Dominic Duggan
Stevens Institute of Technology

1

FINGER TABLES

2


2

Finger Tables

$M = \text{\#nodes}$

$m = \text{\#key bits}$

$M = 2^m$




- Entry i in the finger table of node N is the first node that succeeds or equals $ID_N + 2^i \pmod{M}$
- $FINGER_N(i) = \min \{ID_{N2} \mid ID_{N2} \geq ID_N + 2^i \pmod{M}\}$
- i.e. i^{th} finger points $1/2^{m-i}$ way around the ring

3

3


Example Ring



4

4

Finger Tables




5

5

Example query

- The path of a query for key 54 starting at node 8:



6

6

Node Join and Stabilization

- “Stabilization” protocol contains 6 functions:
 - create()
 - join()
 - stabilize()
 - notify()
 - fix_finger()
 - check_predecessor()

7

7

Node Join – join()

- When node n first starts, it calls $n.join(n')$, for some node n'
- $join()$ function asks n' to find the immediate successor of n

8

8

Node Join – join()


```
// create a new Chord ring.  
n.create()  
    predecessor = nil;  
    successor = n;  
  
// join a Chord ring containing node n'.  
n.join(n')  
    predecessor = nil;  
    successor = n'.find_successor(n);
```

9

9

DHT: Chord Join

- Assume an identifier space [0..7]
- Node n1 joins




10

10

DHT: Chord Join

- Node n2 joins




11

11

DHT: Chord Join

- Node n0 joins

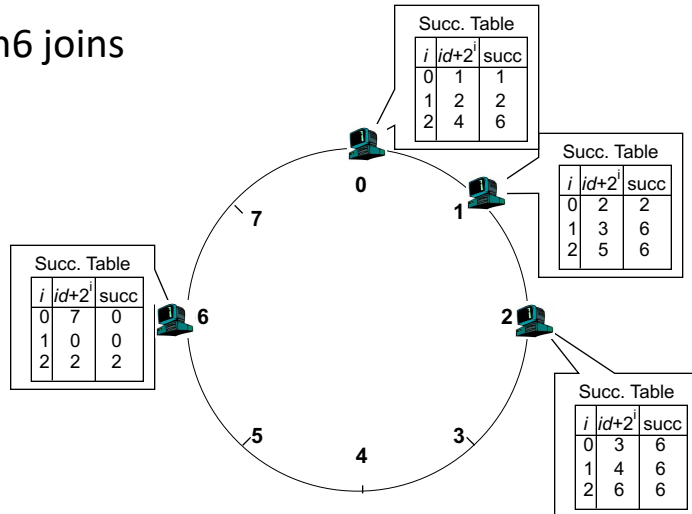


12

12

DHT: Chord Join

- Node n6 joins

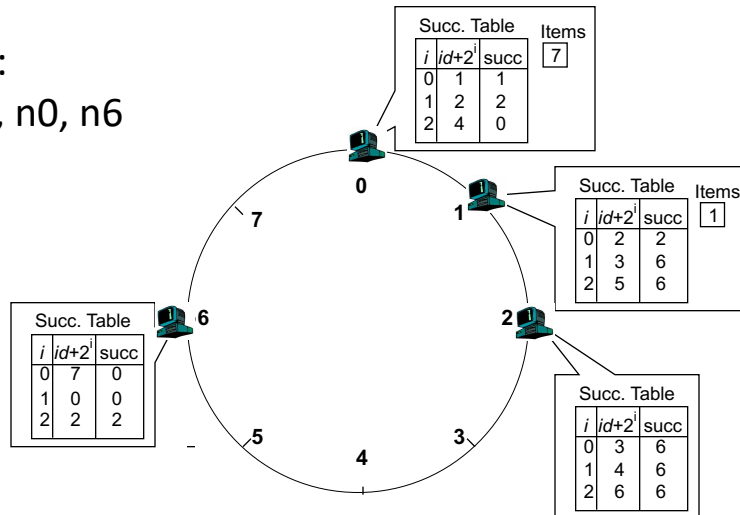


13

DHT: Chord Join

- Nodes:
n1, n2, n0, n6

- Items:
f7, f1




14

14

Scalable Key Location – find_successor()

```
// ask node n to find the successor of id
n.find_successor(id)
  if (id ∈ (n, successor])
    return successor;
  else
    n' = closest_preceding_node(id);
    return n'.find_successor(id);

// search the local table for the highest predecessor of id
n.closest_preceding_node(id)
  for i = m-1 downto 0
    if (finger[i] ∈ (n, id))
      return finger[i];
  return n;
```



15

15

Scalable Key Location – find_successor()

```
// ask node n to find the successor of id
n.find_predecessor(id):
  n' = n
  while (id ∉ (n', n'.successor])
    n' = n'.closest_preceding_finger(id)
  return n'

n.find_successor(id):
  n' = find_predecessor(id)
  return n'.successor

// search the local table for the highest predecessor of id
n.closest_preceding_finger(id)
  for i = m-1 downto 0
    if (finger[i] ∈ (n, id))
      return finger[i];
  return n;
```

16

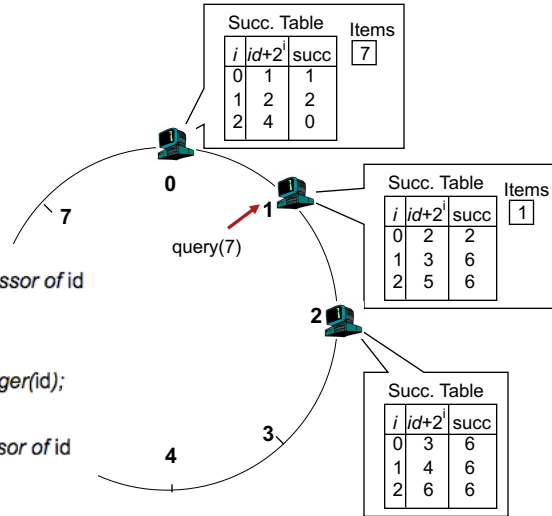
16

DHT: Chord Routing

- Upon receiving a query for item id , a node:
 - Checks whether stores the item locally
 - If not, forwards the query to the largest node in its successor table that does not exceed id

```
// ask node n to find the predecessor of id
n.find_predecessor(id)
n' = n;
while (id > (n', n'.successor])
    n' = n'.closest_preceding_finger(id);
return n'

// ask node n to find the successor of id
n.find_successor(id)
n' = find_predecessor(id);
return n'.successor;
```

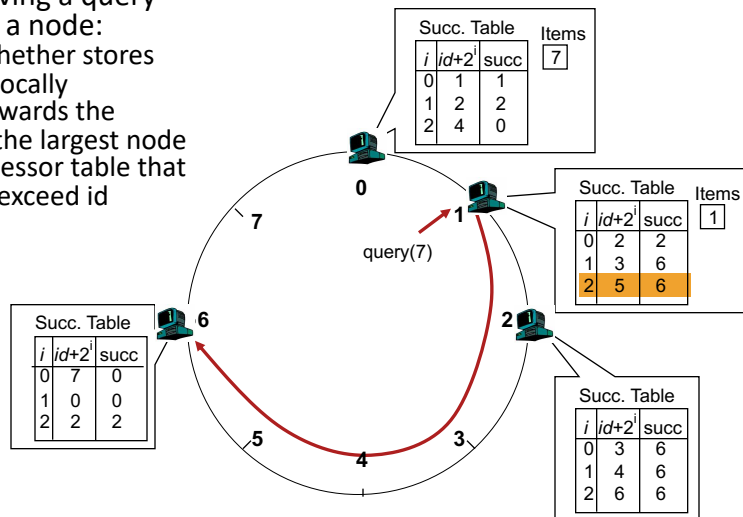


17

17

DHT: Chord Routing

- Upon receiving a query for item id , a node:
 - Checks whether stores the item locally
 - If not, forwards the query to the largest node in its successor table that does not exceed id

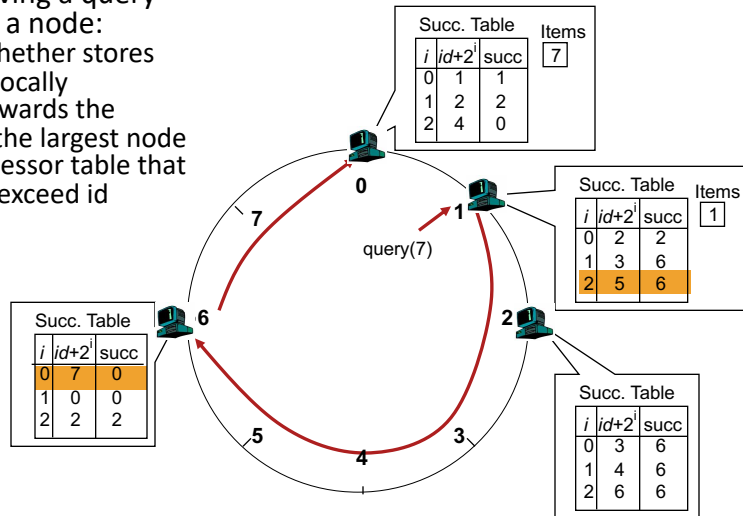


18

18

DHT: Chord Routing

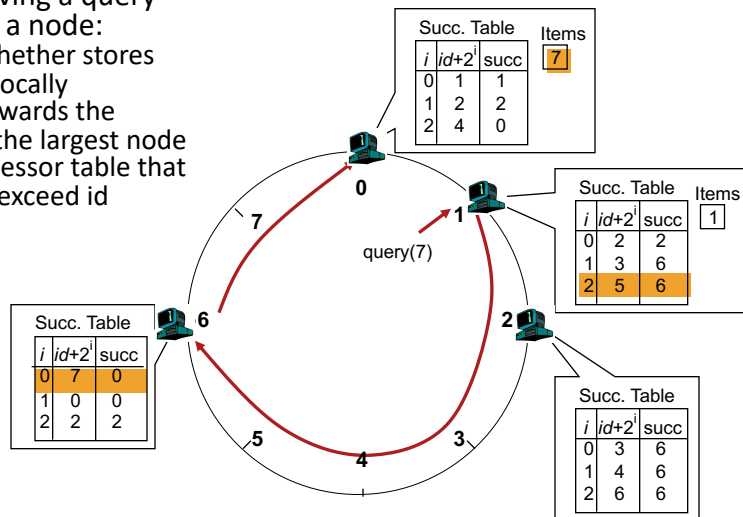
- Upon receiving a query for item id, a node:
 - Checks whether stores the item locally
 - If not, forwards the query to the largest node in its successor table that does not exceed id



19

DHT: Chord Routing

- Upon receiving a query for item id, a node:
 - Checks whether stores the item locally
 - If not, forwards the query to the largest node in its successor table that does not exceed id



20

Chord reliability

- Correct routing table (successors, predecessors, and fingers)
- Primary invariant: correctness of successor pointers
 - Fingers for performance
 - Algorithm is to “get closer” to the target
 - Successor nodes always do this

21

21

Join: Relaxed Approach

- If ring is correct, then routing is correct
- Stabilization
 - Each node periodically runs stabilization routine

22

22

Join: Relaxed Approach

- If ring is correct, then routing is correct
 - Fingers needed for speed only
- Stabilization
 - Each node periodically runs stabilization routine
 - Each node refreshes all fingers by periodically calling `find_successor(N+2i)` for random `i`
 - Periodic cost is $O(\log M)$ per node due to finger refresh

23

23

fix_finger()

```
// called periodically. refreshes finger table entries.
n.fix_finger()
  next = (next + 1) % m ;
  if (next > m-1)
    next = 0;
  finger[next] = find_successor(n + 2next);

// checks whether predecessor has failed.
n.check_predecessor()
  if (predecessor has failed)
    predecessor = nil;
```

24

24

DHT: Chord Summary

- Pros:
 - Guaranteed Lookup
 - $O(\log M)$ per node state and search scope
- Cons:
 - Supporting non-exact match search is hard

25

25

Question

- What assumptions does Chord makes for its search complexity?
- What could change those assumptions?

26