

## SERIALIZABILITY AND RECOVERABILITY

26

26

## Serial Execution


- Run applications one at a time
  - long-running data analytic queries
  - transaction in stored procedure
- Assumptions
  - No I/O (DB entirely in RAM)
  - Single CPU
  - No interactivity
- Ex: Redis

27


27

# Serial Execution

## Interactive transaction:



## Stored procedure:



28

28

# Concurrent Transactions

- Transactions may execute in parallel
  - processor utilization
  - improved average response time
- Schedules
  - Interleaving of update operations of transactions
  - Updates for transaction X must occur in schedule in the same order as in X
- Serial Schedule
  - Ensures isolation of transactions
  - Ensures consistency if each transaction preserves consistency

29

29

## Serial Schedule

- $T_1$   
read A  
 $A := A - 50$   
write A  
read B  
 $B := B + 50$   
write B
- $T_2$   
  
read A  
 $temp := A / 10$   
 $A := A - temp$   
write A  
read B  
 $B := B + temp$   
write B

30

30

## Conflicting Operations

### Read-Write Conflict

- Suppose initial value of A is \$100
- Schedule 1:  
Read A  
Write 150 to A
- Schedule 2:  
Write 150 to A  
Read A

### Write-Write Conflict

- Suppose initial value of A is \$100
- Schedule 1:  
Write 150 to A  
Write 200 to A
- Schedule 2:  
Write 200 to A  
Write 150 to A

31

31

## Conflict Serializability

- Operations O1 and O2 **conflict** if:
  - O1  $\equiv$  read A, O2  $\equiv$  write A
  - O1  $\equiv$  write A, O2  $\equiv$  read A
  - O1  $\equiv$  write A, O2  $\equiv$  write A
- Schedules S and S' are **conflict equivalent** if S' can be obtained from S by swapping non-conflicting operations
- S is **serializable** if S is equivalent to a serial schedule

32

32

## Non-Serial but Serializable Schedule

- |  |  |
|--|--|
| <ul style="list-style-type: none"><li>• T<sub>1</sub><br/>read A<br/>A := A - 50<br/>write A<br/><br/>read B<br/>B := B + 50<br/>write B</li></ul> | <ul style="list-style-type: none"><li>• T<sub>2</sub><br/><br/>read A<br/>temp := A / 10<br/>A := A - temp<br/>write A<br/><br/>read B<br/>B := B + temp<br/>write B</li></ul> |
|--|--|

33

33

## Non-Serial but Serializable Schedule

- $T_1$   
read A  
 $A := A - 50$   
write A  
  
read B  
 $B := B + 50$   
write B
- $T_2$   
  
read A  
 $temp := A / 10$   
 $A := A - temp$   
  
write A  
read B  
 $B := B + temp$   
write B

34

34

## Non-Serial but Serializable Schedule

- $T_1$   
read A  
 $A := A - 50$   
write A  
  
read B  
 $B := B + 50$   
write B
- $T_2$   
  
read A  
  
  
 $temp := A / 10$   
 $A := A - temp$   
write A  
read B  
 $B := B + temp$   
write B

35

35

## Non-Serial but Serializable Schedule

- $T_1$   
read A  
 $A := A - 50$   
write A  
read B  
 $B := B + 50$   
write B
- $T_2$   
  
read A  
 $temp := A / 10$   
 $A := A - temp$   
write A  
read B  
 $B := B + temp$   
write B

36

36

## Non-Serializable Schedule

- $T_1$   
read A  
 $A := A - 50$   
  
write A  
read B  
 $B := B + 50$   
write B
- $T_2$   
  
read A  
 $temp := A / 10$   
 $A := A - temp$   
write A  
read B  
  
 $B := B + temp$   
write B

37

37

## Non-Serializable Schedule

- $T_1$   
read A  
 $A := A - 50$   
  
write A  
read B  
 $B := B + 50$   
  
write B
- $T_2$   
  
read A  
 $temp := A / 10$   
 $A := A - temp$   
write A  
  
read B  
  
 $B := B + temp$   
write B

38

38

## Recoverability

- Successful transaction ends with commit operation
- A schedule is **recoverable** if, whenever  $T_i$  reads data item written by  $T_j$ , commit of  $T_i$  follows commit of  $T_j$

39

39

## Recoverability

- The following schedule is not recoverable:

$T_1$	$T_2$
write A	
	read A
	commit
read B	
commit	

40

40

## Cascading Rollbacks

- The following schedule is recoverable:

$T_1$	$T_2$	$T_3$
read A		
write A		
	read A	
	write A	
		read A
commit		
	commit	
		commit

41

41



## Cascading Rollbacks

- The following schedule is cascadeless:

$T_1$	$T_2$	$T_3$
read A		
write A		
commit		
	read A	
	write A	
	commit	
		read A
		commit

42

42

## Questions

- How could you use locks to ensure serializability?
- How could you use locks to ensure recoverability? Cascadeless schedules?

43

43

## CONCURRENCY CONTROL

44

44

## Implementing Isolation

- Schedules must be
  - serializable
  - recoverable
  - preferably cascadeless
- One-at-a-time execution inefficient
  - Transactions can use the CPU(s) while other transactions are blocked waiting for I/O operations to complete
- Concurrency control

45

45

## Scheduling

- Two operations conflict if they operate on the same data item and at least one of them is a write
  - read-write conflict
  - write-write conflict
- Two read operations can never conflict
- Concurrency control schemes are classified by how they synchronize read and write operations (locking, ordering via timestamps, ...)

46

46

## Two Scheduling Approaches

- Pessimistic - if something can go wrong, it will
  - Operations explicitly synchronized during execution
- Optimistic - in general, nothing will go wrong
  - Synchronization happens at the end of the transaction

47

47

## Lost Updates

- $T_1$   
read A  
 $A := A - 50$   
  
write A
- $T_2$   
  
read A  
 $A := A + 100$   
  
write A

48

48

## Locking for Isolation

- $T_1$   
lock-X A  
read A  
 $A := A - 50$   
write A  
unlock A
- $T_2$   
  
lock-X A  
read A  
 $A := A + 100$   
write A  
unlock A

49

49

## Examples of lock coverage

- We could have one lock per object
- ... or one lock for the whole database
- ... or one lock for a category of objects
  - Tree
  - Table, row, column
- All transactions must use the same rules!
- “Write” locks for updates

50

50

## Lock-Based Concurrency Control

- Two modes of locking:
  - shared (read-only)
  - exclusive (read-write)
- Lock acquisition
  - transaction blocks if lock not available
  - update shared lock to exclusive
  - downgrade exclusive lock to shared
- Possibility of **deadlock**

51

51

## Lock-Based Concurrency Control

- Lock Compatibility

Request \ Held	S	X
by $T_i$ \ by $T_j$		
S	true	false
X	false	false

- Lock manager implemented using readers-writers algorithm

52

52

## Locking for Isolation

- $T_1$

```
lock-X A; read A
A := A - 50
write A; unlock A
```

```
lock-X B; read B
B := B + 50
write B; unlock B
```

- $T_2$

```
lock-X B; read B
B := B - 100
write B; unlock B
```

```
lock-X A; read A
A := A + 100
write A; unlock A
```

53

53

## Locking for Serializability

- $T_1$   
lock-X A; lock-X B  
read A;  $A := A - 50$   
write A  
read B;  $B := B + 50$   
write B  
unlock A; unlock B
- $T_2$   
lock-X B; lock-X A  
read B;  $B := B - 100$   
write B  
read A;  $A := A + 100$   
write A  
unlock B; unlock A

54

54

## Locking for Serializability

- $T_1$   
lock-X A; read A  
 $A := A - 50$   
write A;  
lock-X B; read B  
 $B := B + 50$   
write B;  
unlock A; unlock B
- $T_2$   
lock-X B; read B  
 $B := B - 100$   
write B;  
Lock-X A; read A  
 $A := A + 100$   
write A;  
unlock B; unlock A

55

55

## Two-Phase Locking Protocol (2PL)

- Protocol:
  - Phase 1 (Growing)
    - transaction may obtain or upgrade locks
    - transaction may not release or downgrade locks
  - Phase 2 (Shrinking)
    - transaction may release or downgrade locks
    - transaction may not obtain or upgrade locks
- 2PL ensures serializability
  - Prevent cycles in temporal dependencies
- 2PL increases chance of deadlock

56

56

## Locking for Recoverability

- |   |  |
|---|--|
| <ul style="list-style-type: none"><li>• <math>T_1</math><br/>lock-X A; read A<br/><math>A := A - 50</math>; write A<br/>lock-X B; read B<br/><math>B := B + 50</math>; write B<br/>unlock A; unlock B<br/><br/>commit</li></ul> | <ul style="list-style-type: none"><li>• <math>T_2</math><br/><br/><br/><br/><br/>lock-X B; read B<br/><math>B := B - 100</math>; write B<br/>lock-X A; read A<br/><math>A := A + 50</math>; write A<br/>unlock B; unlock A<br/><br/>commit</li></ul> |
|---|--|

57

57



## Locking for Recoverability

- T1
  - lock-X A
  - read A; A := A - 50
  - write A;
  - lock-X B
  - read B; B := B + 50
  - write B
  - commit; unlock A; unlock B
- T2
  - lock-X B
  - read B; B := B - 100
  - write B
  - lock-X A
  - read A; A := A + 50
  - write A
  - commit; unlock B; unlock A

58

58

## Variations on 2PL

- Strict 2PL
  - Transaction must hold all **exclusive** locks until it commits or aborts
  - Ensures cascadeless schedules
- Rigorous 2PL
  - Transaction must hold **all** locks until it commits or aborts
  - Ensures transactions can be serialized in the order in which they commit

59

59