

STATE MACHINE: ORDERING UPDATES

99

99


State Machine Replication

- Active replication (State machine)
 - Peer-to-peer replicas
 - Each replica is **deterministic** state machine
 - **Operations** executed in same order on all replicas
 - All updates are totally ordered

100


100

Consistent Ordering of Updates



101

State Machine Replication



- How to maintain a single order in the face of concurrent client requests?

102

102

Versions of Replicated Data

- Replicated data items have “versions”
 - I.e. can't just say “ $X_p=3$ ”.
 - X_p has *timestamp* $[7,q]$
 - X_p has *value* 3
 - Timestamp
 - must increase monotonically
 - includes a process id to break ties

103

103

Read

- Wait until Q_R processes reply
- Use value with largest timestamp
 - Break ties by looking at the pid
 - For example
 - $[6,x] < [9,a]$
 - $[7,p] < [7,q]$
 - *Even if a process owns a replica, it can't just trust its own data*

104

104

Write

- Can't support incremental updates
 - $x = x + 1$
 - Insert into a queue
- Quorum
 - Use a commit protocol
- How to determine the version number
 - Voting protocol

105

105

Protocol

1. Propose the write: "I would like to set $X=3$ "
2. Members "lock" the variable against reads, *put the request into a queue of pending writes*, and send back:
"I propose time [t,pid]"
Time is a logical clock.
3. Initiator collects replies, hoping to receive Q_w

$\geq Q_w$ OKs

$< Q_w$ OKs

Compute maximum of
proposed [t,pid] pairs.
Commit at that time

Abort

106

106

Voting Based on Logical Time

- Logical clocks
 - See mutual exclusion algorithm with logical time
- Update source takes the maximum
 - Commit message: “commit at [t,pid]”
 - Group member: if vote considered:
 - deliver committed updates in timestamp order
 - Group members: if vote not considered:
 - discard the update

107


107

Where are the updates?


- Each member: queue of uncommitted updates
 - Survives crash and restart
- Example: Process p
 - (u_2 : [1,p] pending), (u_1 : [2,p] pending)
 - Neither can be delivered

108


108




109



110



111



112

What if “my vote wasn’t used?”

- Process
 - had a pending update
 - discovers it wasn’t used
- Discard the request
 - Otherwise block forever (why?)
 - Ignoring the request won’t hurt (why?)

113

113

Which votes got counted?

- Need to know which votes were “counted”
 - E.g. suppose A,B,C,D,E and they vote:
 - {[17,A] [19,B] [20,C] [200,D] [21,E]}
 - Vote from D is lost
 - the maximum is picked as [21,E]
 - Remember that the votes used to make this decision were from {A,B,C,E}

114

114

Recovery

- First recover queue of pending updates
- Next, learn the outcome of the operation
 - Contact Q_R other replicas
- Check if own vote counted (if committed)
 - If so, apply update
 - If not, discard update

115

115

Read requests while updates pending...

- Suppose a read while updates pending
 - Wait until those commit, abort, or are discarded
 - Otherwise process might not see its own updated value

116

116

Why is this “safe”?

- Commit: only move pending update to later
 - Discard pending update if vote not counted
 - Result: inconsistent replica
 - *but we always look at Q_R replicas*
 - Why we can't support incremental operations (insert, etc)

117

117

Why is this “safe”?

- Commit: moves pending update to front of Q
- Once a committed update reaches front of Q:
 - ...no update can be committed at earlier time!
- Any “future” update gets later time

118

118

Why this works

- Everyone uses same commit time for an update
 - Can't deliver update unless [t,pid] is smallest
 - and is committed
 - Hence updates in same order at all replicas

119

119

Observations

- The protocol requires many messages
 - Could use IP multicast for first and last round
 - Need reliability
- Commit messages must be reliably delivered
 - Otherwise uncommitted updates on front of Q...
- 2PC and 3PC may block (why?)
 - FLP: *any* quorum write protocol can block

120

120