

## **VIEWSTAMP REPLICATION: ORDERING UPDATES**

44

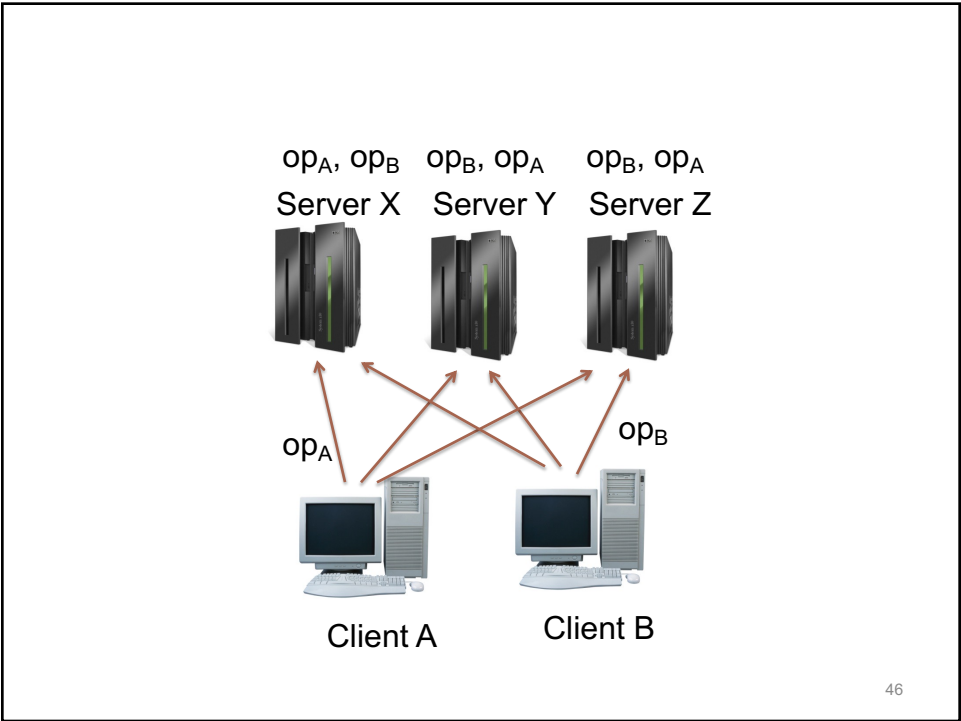
44

## **Quorum Consensus**

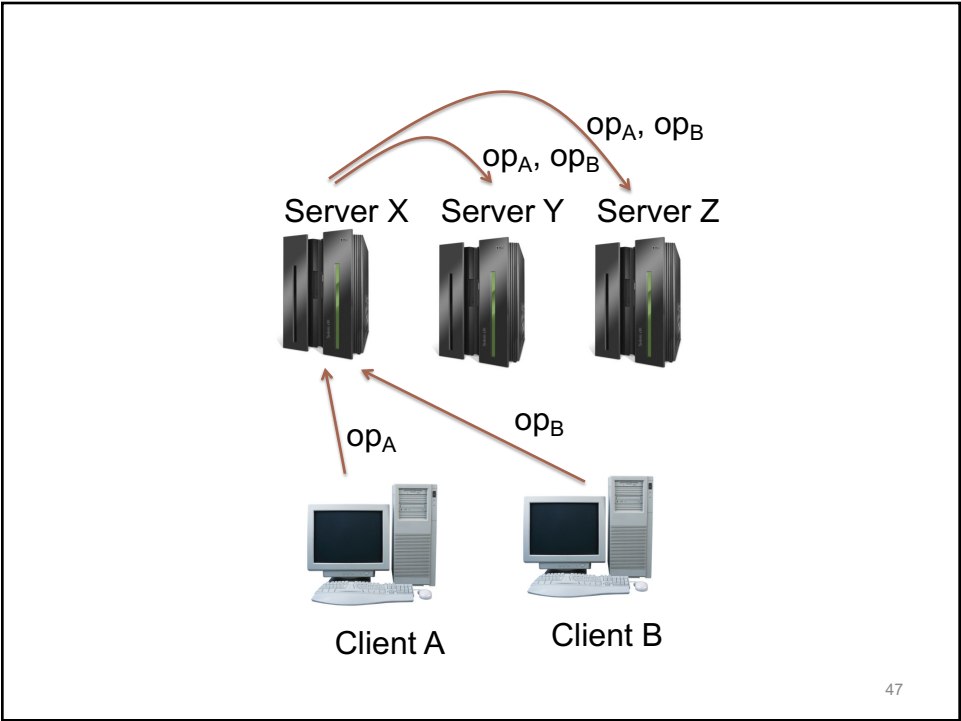
- Crash-stop failures
- Requires  $2F+1$  replicas
  - Operations must intersect for at least one replica
  - Want availability for both reads and writes
  - Read and write quorums of  $F+1$  nodes

45

45



46



47

## Viewstamp Replication

- Primary-backup
- System moves through a sequence of views
  - Primary runs the protocol
  - Replicas do a view change if it fails

48

48

## Replica state


- A **replica id**  $i$  (between 0 and  $N-1$ )
  - Replica 0, replica 1, ...
- A **view number**  $v\#$ , initially 0
- **Primary** is the replica with id  
 $i = v\# \bmod N$
- A **log** of  $\langle op, op\#, status \rangle$  entries
  - Status = **prepared** or **committed**

49

49

Client knows current view #

View # = 2  $\Rightarrow$  primary pid = 2 (i.e., r)




50

50

Primary chooses logical timestamp for update, e.g.  $LT(u_1)=9$


Op # = (View #, time)



51

51

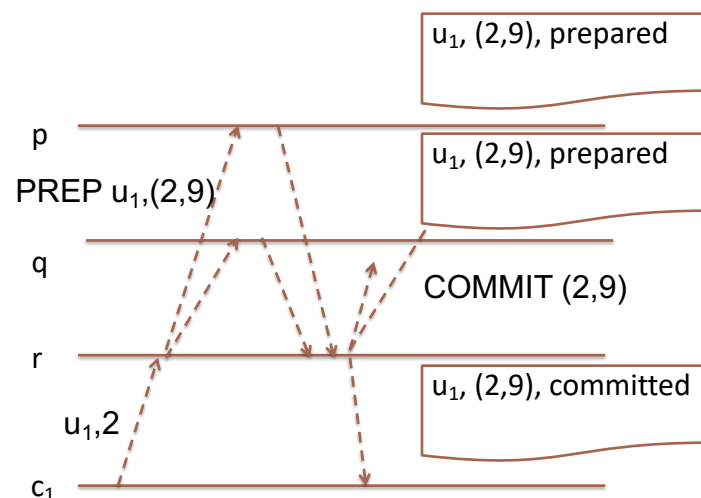
Primary waits for  $\geq f$  replicas to respond ( $\geq f+1$  total)



52

52

Client notified immediately after acks



53

53

## **VIEWSTAMP REPLICATION: VIEW CHANGE**

54

54

## **View Changes**

- Used to mask primary failures
- Replicas monitor the primary
- Replica requests next primary to do a view change

55

55


## Correctness Requirement

- Operation order must be preserved by a view change
- For operations that are **visible**
  - executed by server
  - client received result
- An operation can be visible if it prepared at  $f+1$  replicas
  - this is the **commit point**

56

56


Primary  $r$  has crashed after assigning  $LT(u_1)=9$



57

57

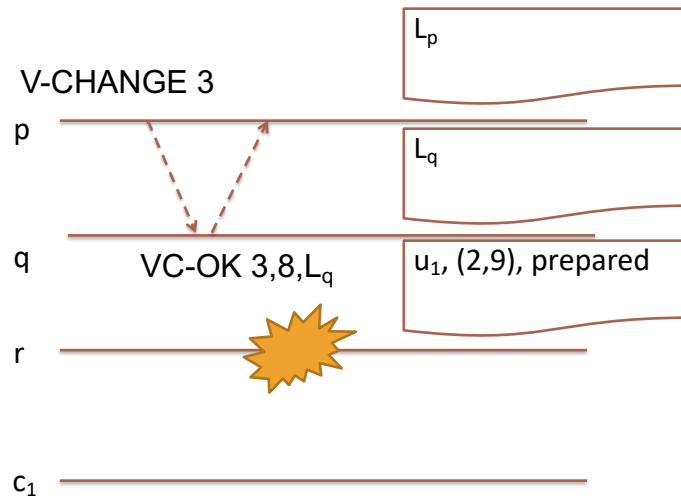
View # increments (to 3), new primary is p, start view change



58

58

Replica q acks view change, returns last op# it knew of and its log ( $L_q$  contains any operations that p never heard of)




59

59



Primary p lets replicas know (via its log  $L_p'$ ) of ops that replicas were never informed of before old primary crashed



60

60


## View Change

- New primary may not know of all updates from old primary
- New primary asks replicas for their logs
- Any committed operation was acked by a quorum, so must be in log of a surviving replica
  - Primary takes the max of the logs returned
  - That log has most recent updates

61

Second client  $c_2$  knows current view #


View# = 3  $\Rightarrow$  primary pid = 0 (i.e., p)



62

62


Primary p gets the other replicas to prepare to commit  $u_2$



63

63

Viewstamp (3,9) avoids confusion between  $u_1$  and  $u_2$  at  $r$   
 (Replica  $r$  will not think  $p$  is committing  $u_1$ )



64

64

## Persistent State

- Voting protocol: votes must survive failure
  - Save queue of pending updates on disk
- Viewstamp: primary can respond to client without recording commit on disk
  - View change: recover commitment from logs of surviving replicas
- Only need to persist state after view change
  - So if we crash and recover, we know view# when we crashed
  - Even that unnecessary with more expensive recovery protocol

65

## Summary

- Primary-Backup
  - Order updates at the primary
  - Order preserved by view change
- Quorum Consensus
- Split brain?
- What about FLP?

66