PAXOS

67

67

RSM: Split-Brain

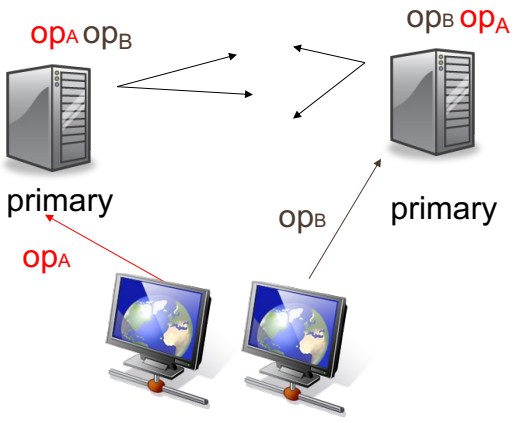

68

68

Paxos: Consensus Box

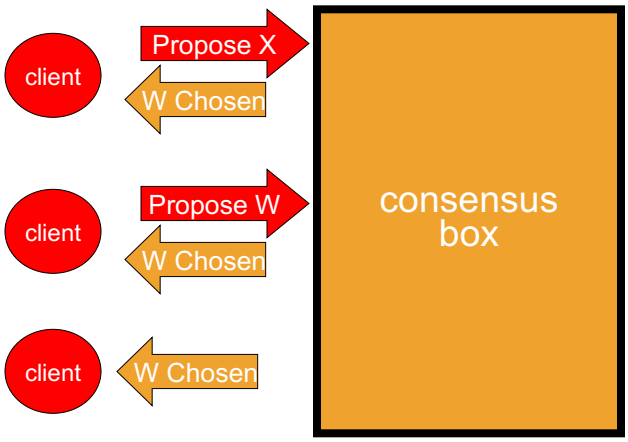

- Collects proposed values
- Picks one proposed value
- Remembers it forever

69

69

Paxos: Normal Execution

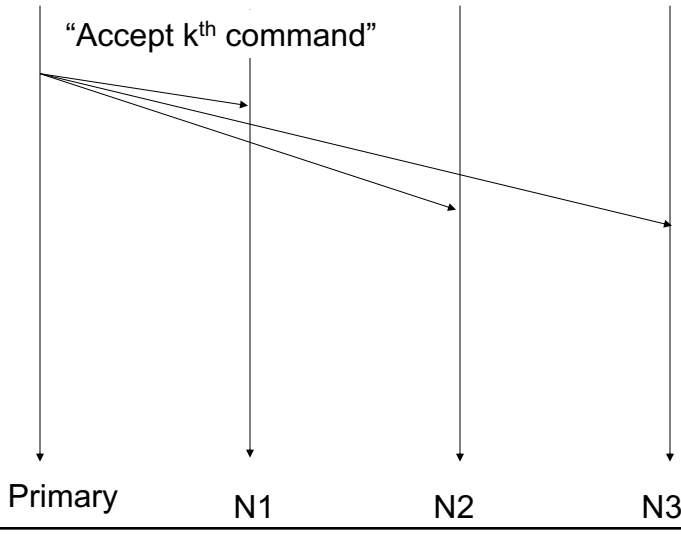

70

70

Paxos: Normal Execution

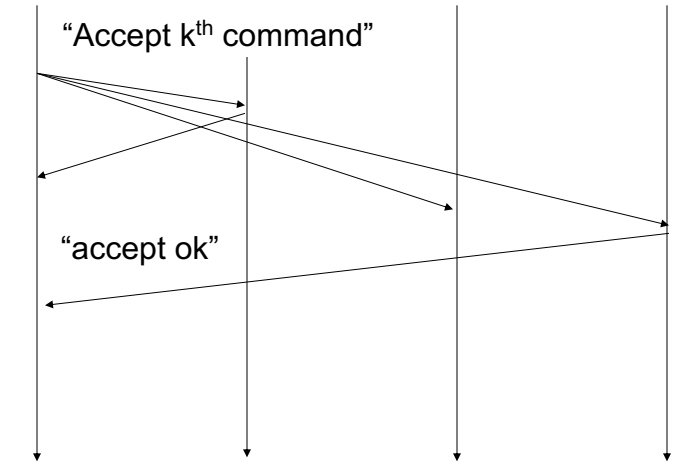

71

71

Paxos: Normal Execution

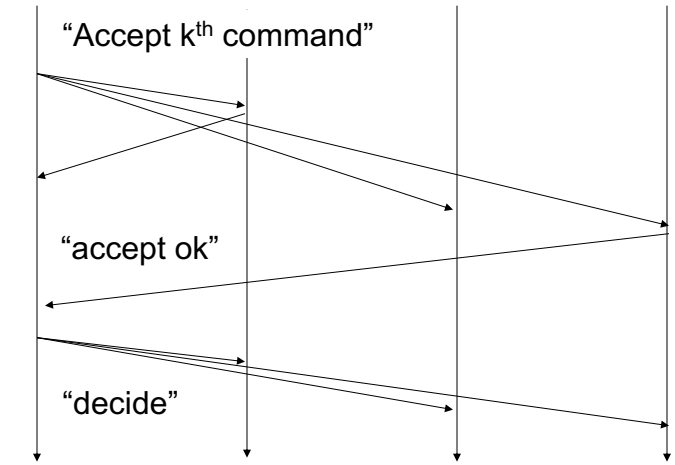

72

72

Paxos: Split Brain

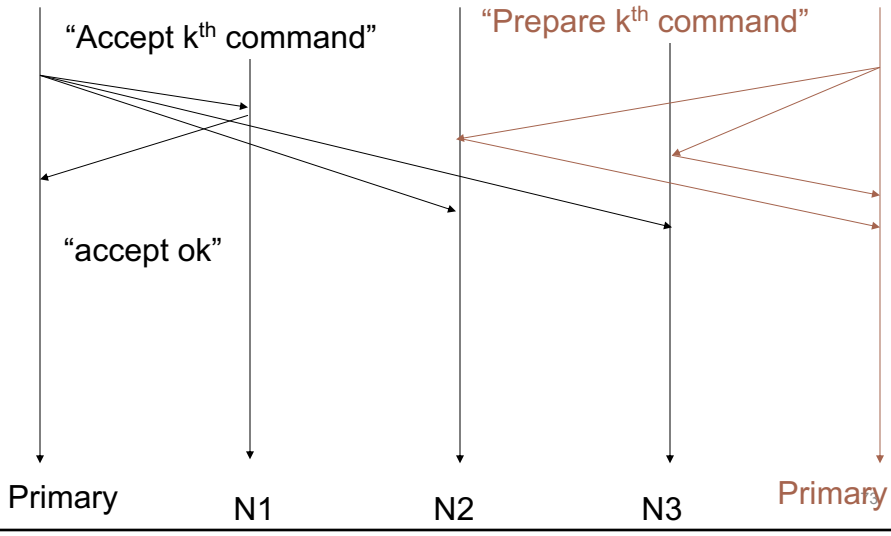

73

Paxos: Split Brain

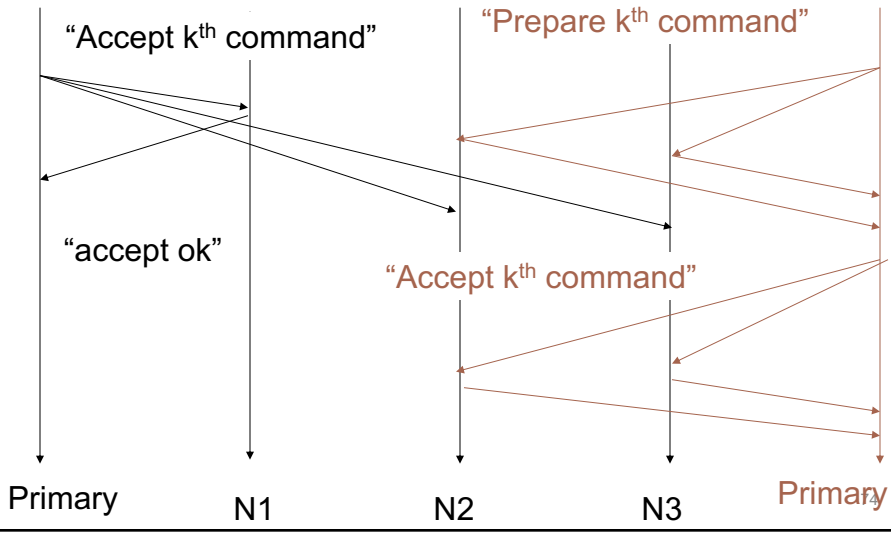

74

Paxos: general approach

- One (or more) node decides to be the leader
- Leader proposes a value and solicits acceptance from others (acceptors)
- Leader announces result **or tries again**

75

75

Paxos requirement

- Correctness (safety):
  - All nodes agree on the same value
  - The agreed value X has been proposed by some node
- Fault-tolerance:
  - If less than N/2 nodes fail, the rest nodes should reach agreement *eventually w.h.p*
  - Liveness is not *guaranteed*

76

76

Why is agreement hard?

- What if >1 nodes become leaders simultaneously?
- What if there is a network partition?
- What if a leader crashes in the middle of solicitation?
- What if a leader crashes after deciding but before announcing results?
- What if the new leader proposes different values than already decided value?

77

77

Paxos setup

- Each node runs as a *proposer*, *acceptor* and *learner*
- Proposer (leader) proposes a value and solicit acceptance from acceptors
- Leader announces the chosen value to learners

78

78

Paxos setup

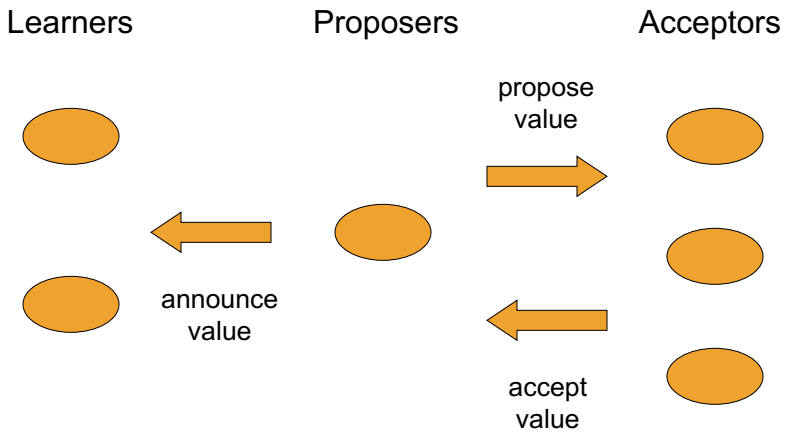

79

79

Strawman 1: Single Acceptor

- Designate a single node X as acceptor (e.g. one with smallest id)
  - Each *proposer* sends its value to X
  - X decides on one of the values
  - X announces its decision to all *learners*
- Problem?
  - Failure of the single acceptor halts decision
  - Need multiple acceptors!

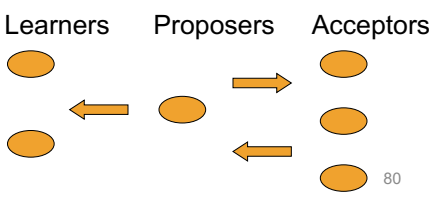

80

80

Strawman 2: multiple acceptors

- Each proposer (leader) proposes to all acceptors
- Each acceptor accepts the first proposal it receives and rejects other proposals
- If the leader receives positive replies from a majority of acceptors, it chooses its own value
  - There is at most 1 majority, hence only a single value is chosen
- Leader sends chosen value to all learners

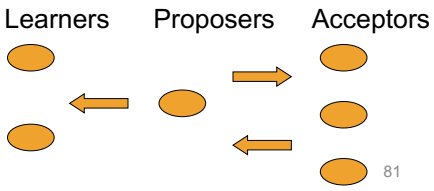

81

81

Strawman 2: multiple acceptors

- Each proposer (leader) proposes to all acceptors
- Each acceptor accepts the first proposal it receives
  - Rejects other proposals
- If the leader receives positive replies from a majority of acceptors, it chooses its own value
  - There is at most 1 majority
- Leader sends chosen value to all learners
- Problem:
  - What if multiple leaders propose simultaneously so there is no majority accepting?

82

82

Paxos solution

- Proposals (for a value e.g. k^{th} command) are ordered by proposal #
- Each acceptor must accept the first proposal that it receives
- Each acceptor may accept multiple proposals
  - If a proposal with value v is chosen, all higher proposals chosen have value v

83

83

Paxos solution

- Proposals (for a value e.g. k^{th} command) are ordered by proposal #
- Each acceptor must accept the first proposal that it receives
- Each acceptor may accept multiple proposals
  - If a proposal with value v is chosen, all higher proposals chosen have value v
  - If a proposal with value v is chosen, all higher proposals accepted by any acceptor have value v

84

84

Paxos solution

- Proposals (for a value e.g. k^{th} command) are ordered by proposal #
- Each acceptor must accept the first proposal that it receives
- Each acceptor may accept multiple proposals
  - If a proposal with value v is chosen, all higher proposals chosen have value v
  - If a proposal with value v is chosen, all higher proposals accepted by any acceptor have value v
  - If a proposal with value v is chosen, all higher proposals issued by any proposer have value v

85

85

Paxos solution

- Proposals (for a value e.g. k^{th} command) are ordered by proposal #
- Each acceptor must accept the first proposal that it receives
- Each acc
  - If a pro
  - If a pro
  - If a pro

Before proposing value v for proposal n, proposer will poll acceptors for
• Promise that they will not accept any future proposals < n
• What value if any that they accepted for highest numbered proposal < n

by any proposer have value v

86

86

Paxos operation: node state

- Each node maintains:
  - n_a, v_a: highest proposal # and its corresponding accepted value
    - initially null
  - n_h: highest proposal # seen
  - my_n: my proposal # in current Paxos

87

87

Paxos operation: 3P protocol

- Phase 1 (Prepare)
  - A node decides to be leader (and propose)
  - Leader chooses $my_n > n_h$
  - Leader sends <prepare, my_n> to all nodes

88

88

Paxos operation: 3P protocol

- Phase 1 (Prepare)
  - A node decides to be leader (and propose)
  - Leader chooses my$n$ > n_h
  - Leader sends <prepare, my$n$> to all nodes
  - Upon receiving <prepare, n>

    If n < n_h

    reply <prepare-reject>

    else /* n > n_h */

    n_h = n

    reply <prepare-ok, n_a,v_a>

89

89

Paxos operation: 3P protocol

- Phase 1 (Prepare)
  - A node decides to be leader (and propose)
  - Leader chooses my$n$ > n_h
  - Leader sends <prepare, my$n$> to all nodes
  - Upon receiving <prepare, n>

    If n < n_h

    reply <prepare-reject>

    else /* n > n_h */

    n_h = n

    reply <prepare-ok, n_a,v_a>

This node will not accept any proposal lower than n

90

90

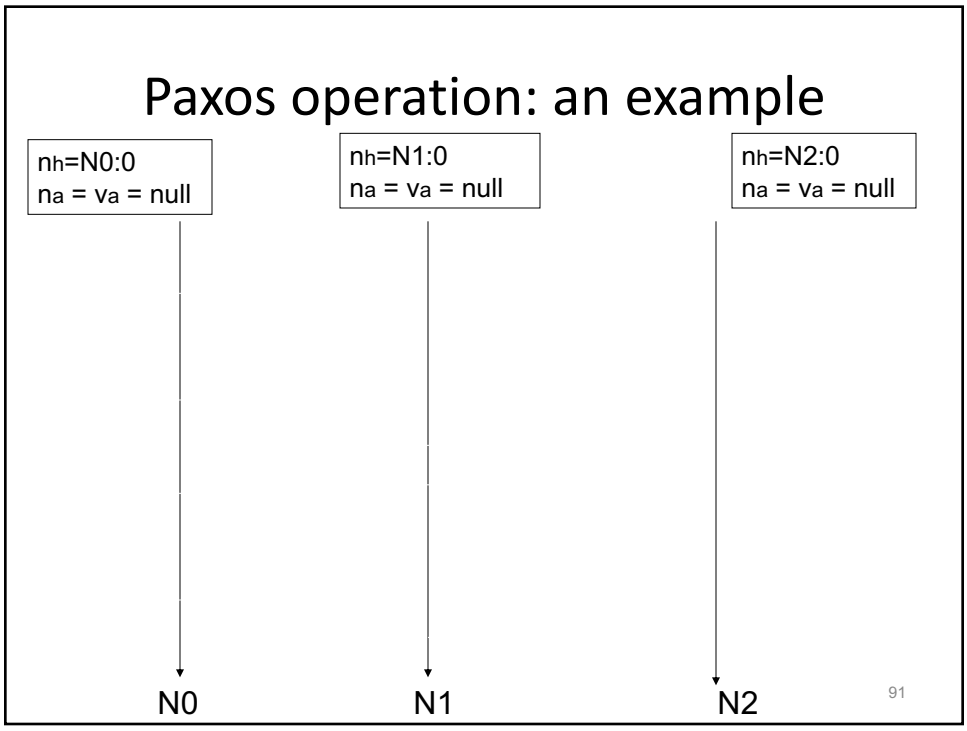
Paxos operation: an example
nh=N0:0
na = va = null
nh=N1:0
na = va = null
nh=N2:0
na = va = null
N0
N1
N2
91

91

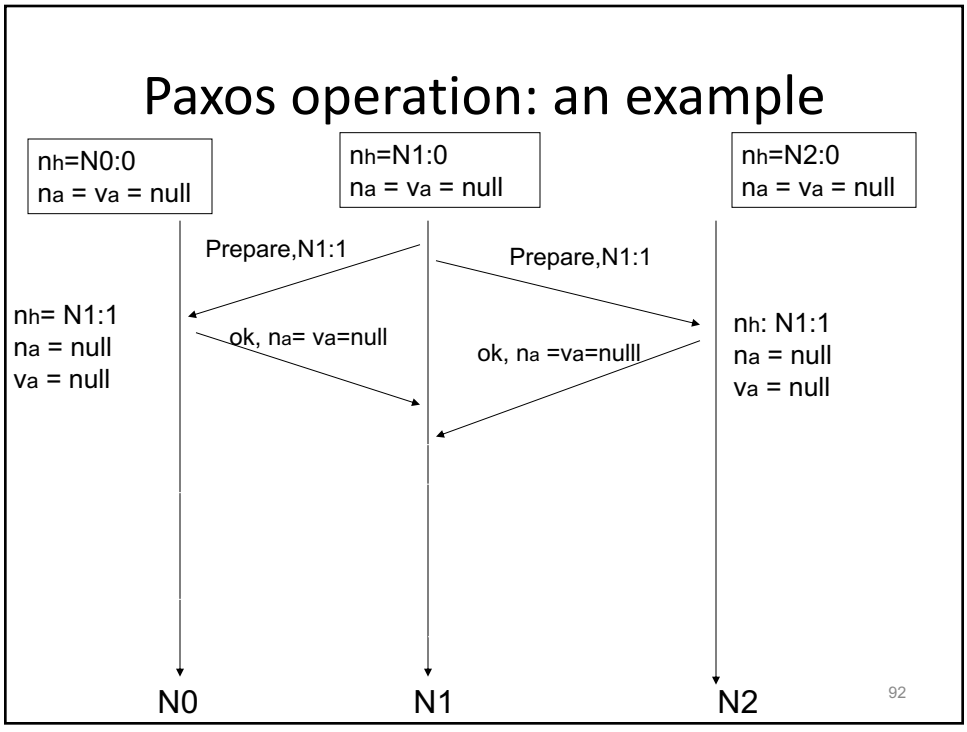
Paxos operation: an example
nh=N0:0
na = va = null
nh=N1:0
na = va = null
nh=N2:0
na = va = null
Prepare,N1:1
Prepare,N1:1
nh= N1:1
na = null
va = null
ok, na= va=null
ok, na =va=nulll
nh: N1:1
na = null
va = null
N0
N1
N2
92

92

Paxos operation

- Phase 2 (Accept):
  - If leader gets prepare-ok from a majority
    V = non-empty value corresponding to the highest n_a received
    If V = null, then leader can pick any V
    Send <accept, my_n, V> to all nodes

Paxos operation

- Phase 2 (Accept):
  - If leader gets prepare-ok from a majority
    V = non-empty value corresponding to the highest n_a received
    If V = null, then leader can pick any V
    Send <accept, my_n, V> to all nodes
  - If leader fails to get majority prepare-ok
    - Delay and restart Paxos

Paxos operation

- Phase 2 (Accept):
  - If leader gets prepare-ok from a majority
    V = non-empty value corresponding to the highest n_a received
    If V = null, then leader can pick any V
    Send <accept, my_n, V> to all nodes
  - If leader fails to get majority prepare-ok
    - Delay and restart Paxos
  - Upon receiving <accept, n, V>
    If $n < n_h$
      reply with <accept-reject>
    else
      $n_a = n$; $v_a = V$; $n_h = n$
      reply with <accept-ok>

95

95

Paxos operation: an example

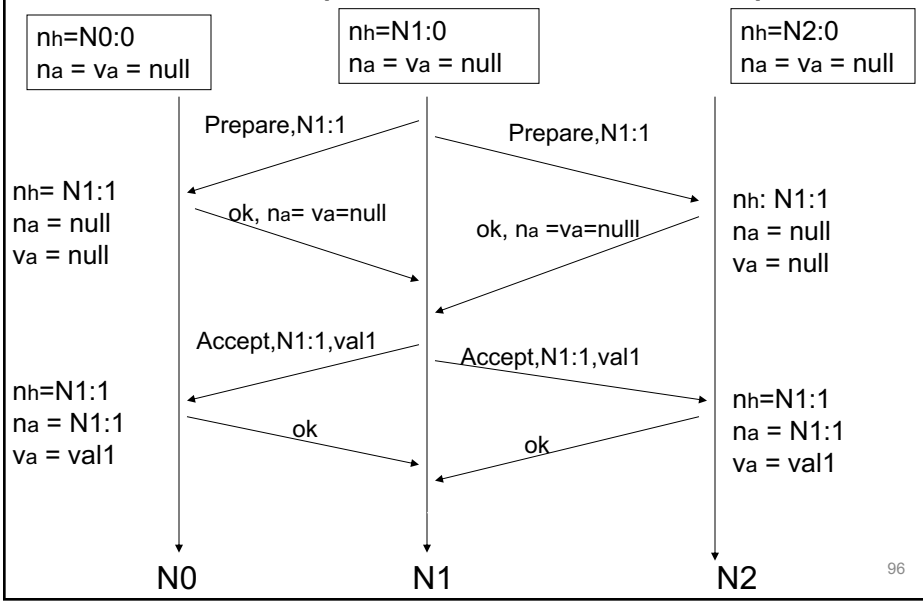

96

96

Paxos operation

- Phase 3 (Decide)
  - If leader gets accept-ok from a majority
    - Send <decide, v_a> to all nodes

Paxos operation

- Phase 3 (Decide)
  - If leader gets accept-ok from a majority
    - Send <decide, v_a> to all nodes
  - If leader fails to get accept-ok from a majority
    - Delay and restart Paxos

Paxos operation: an example

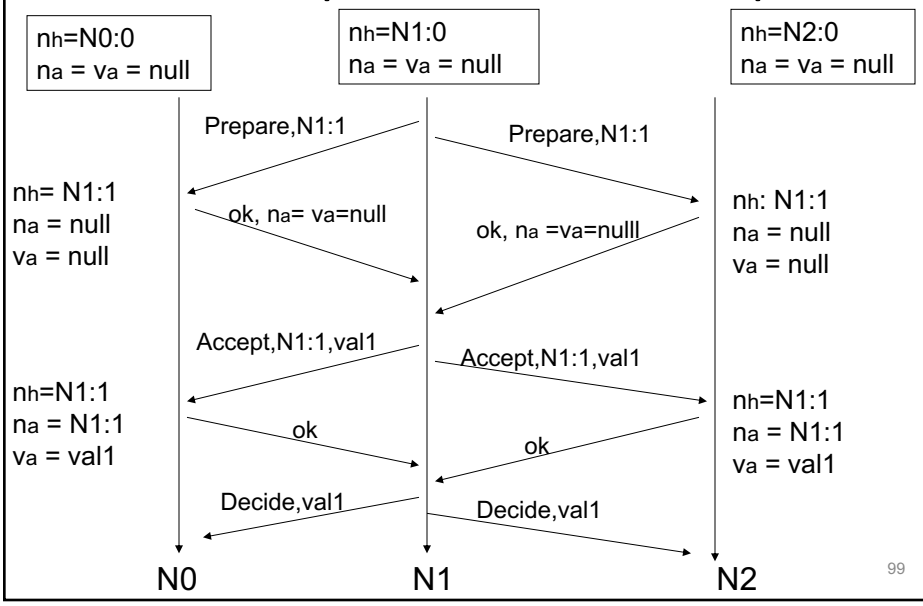

99

99

Paxos properties

- When is the value V chosen?
  1. When leader receives a majority prepare-ok and proposes V
  2. When a majority of nodes accept V
  3. When the leader receives a majority accept-ok for value V

100

100

Understanding Paxos

- What if more than one leader is active?
- Suppose two leaders use different proposal number, N0:10, N1:11
- Can both leaders see a majority of prepare-ok?

Understanding Paxos

- What if leader fails while sending accept?
- What if a node fails after receiving accept?
  - If it doesn't restart …
  - If it reboots …
- What if a node fails after sending prepare-ok?
  - If it reboots …